



# INNOVATE2018

## ONLINE CONFERENCE



# Build Business Outcomes with a Modern Data Architecture (Level 200)

Craig Stires, Head of Big Data & Analytics

# The old approaches aren't good enough

# The old approaches aren't good enough

## People



Too much specialization  
required and other heavy  
lifting keeps teams stagnant

# The old approaches aren't good enough

## People



Too much specialization required and other heavy lifting keeps teams stagnant

## Process



Long procurement and deployment cycles create missed opportunities

# The old approaches aren't good enough

## People



Too much specialization required and other heavy lifting keeps teams stagnant

## Process



Long procurement and deployment cycles create missed opportunities

## Technology



Legacy systems are inadequate or too expensive to scale for evolving use cases

# Why customers build data platforms on AWS

# Why customers build data platforms on AWS

## People



Infrastructure as code allows  
for fast iterations and  
auditable changes, at the  
speed of business

# Why customers build data platforms on AWS

## People



Infrastructure as code allows for fast iterations and auditable changes, at the speed of business

## Process



Some services are instantly available, others take minutes to provision, not months



# Why customers build data platforms on AWS

## People



Infrastructure as code allows for fast iterations and auditable changes, at the speed of business

## Process



Some services are instantly available, others take minutes to provision, not months

## Technology



AWS has the broadest and deepest cloud platform, and is built for security and governance

# Let's go on a common customer journey

# Meet EarEcstasy, as they move from B2B to B2C

## 2005

EarEcstasy manufactures headsets. They ran a traditional B2B business since 2005, selling through distribution and retail channels.



\* This case is representative of a common customer journey, but EarEcstasy isn't an actual business

# Meet EarEcstasy, as they move from B2B to B2C

## 2005

EarEcstasy manufacturers headsets. They ran a traditional B2B business since 2005, selling through distribution and retail channels.



## 2018

In 2018, they launched their first “Smart Buds”. These wireless headsets have voice enablement, GPS tracking, and heartrate monitors built in, and the device syncs with the users mobile phone via Bluetooth. The mobile app also supports scene detection.



\* This case is representative of a common customer journey, but EarEcstasy isn't an actual business

# EarEcstasy needs to answer new questions and move faster



Lim, Head of Finance



Raymond, Head of Product

# EarEcstasy needs to answer new questions and move faster

How do **quality issues** impact cost of production?

Can I look at **supplier performance** over time?

How can we reduce our **inventory holding costs**?



Lim, Head of Finance



Raymond, Head of Product

# EarEcstasy needs to answer new questions and move faster

How do **quality issues** impact cost of production?

Can I look at **supplier performance** over time?

How can we reduce our **inventory holding costs**?



Lim, Head of Finance

**Which regions** are the new earbuds selling well?

What is the **demand forecast** by product category?

What is the **social sentiment** about our products?




Raymond, Head of Product

# To answer new questions quickly, we look to a modern data architecture design

Massive upfront costs  
Overprovisioned capacity  
Long implementation times



Pay as you go, for what you use  
Decoupled pipelines and engines  
Experimentation platform





# Outcome 1: Modernize and consolidate

# Start with a set of specific questions to answer, then work backwards to the data required


How do **quality issues** impact cost of production?

Can I look at **supplier performance** over time?


How can we reduce our **inventory holding costs**?



Lim, Head of Finance



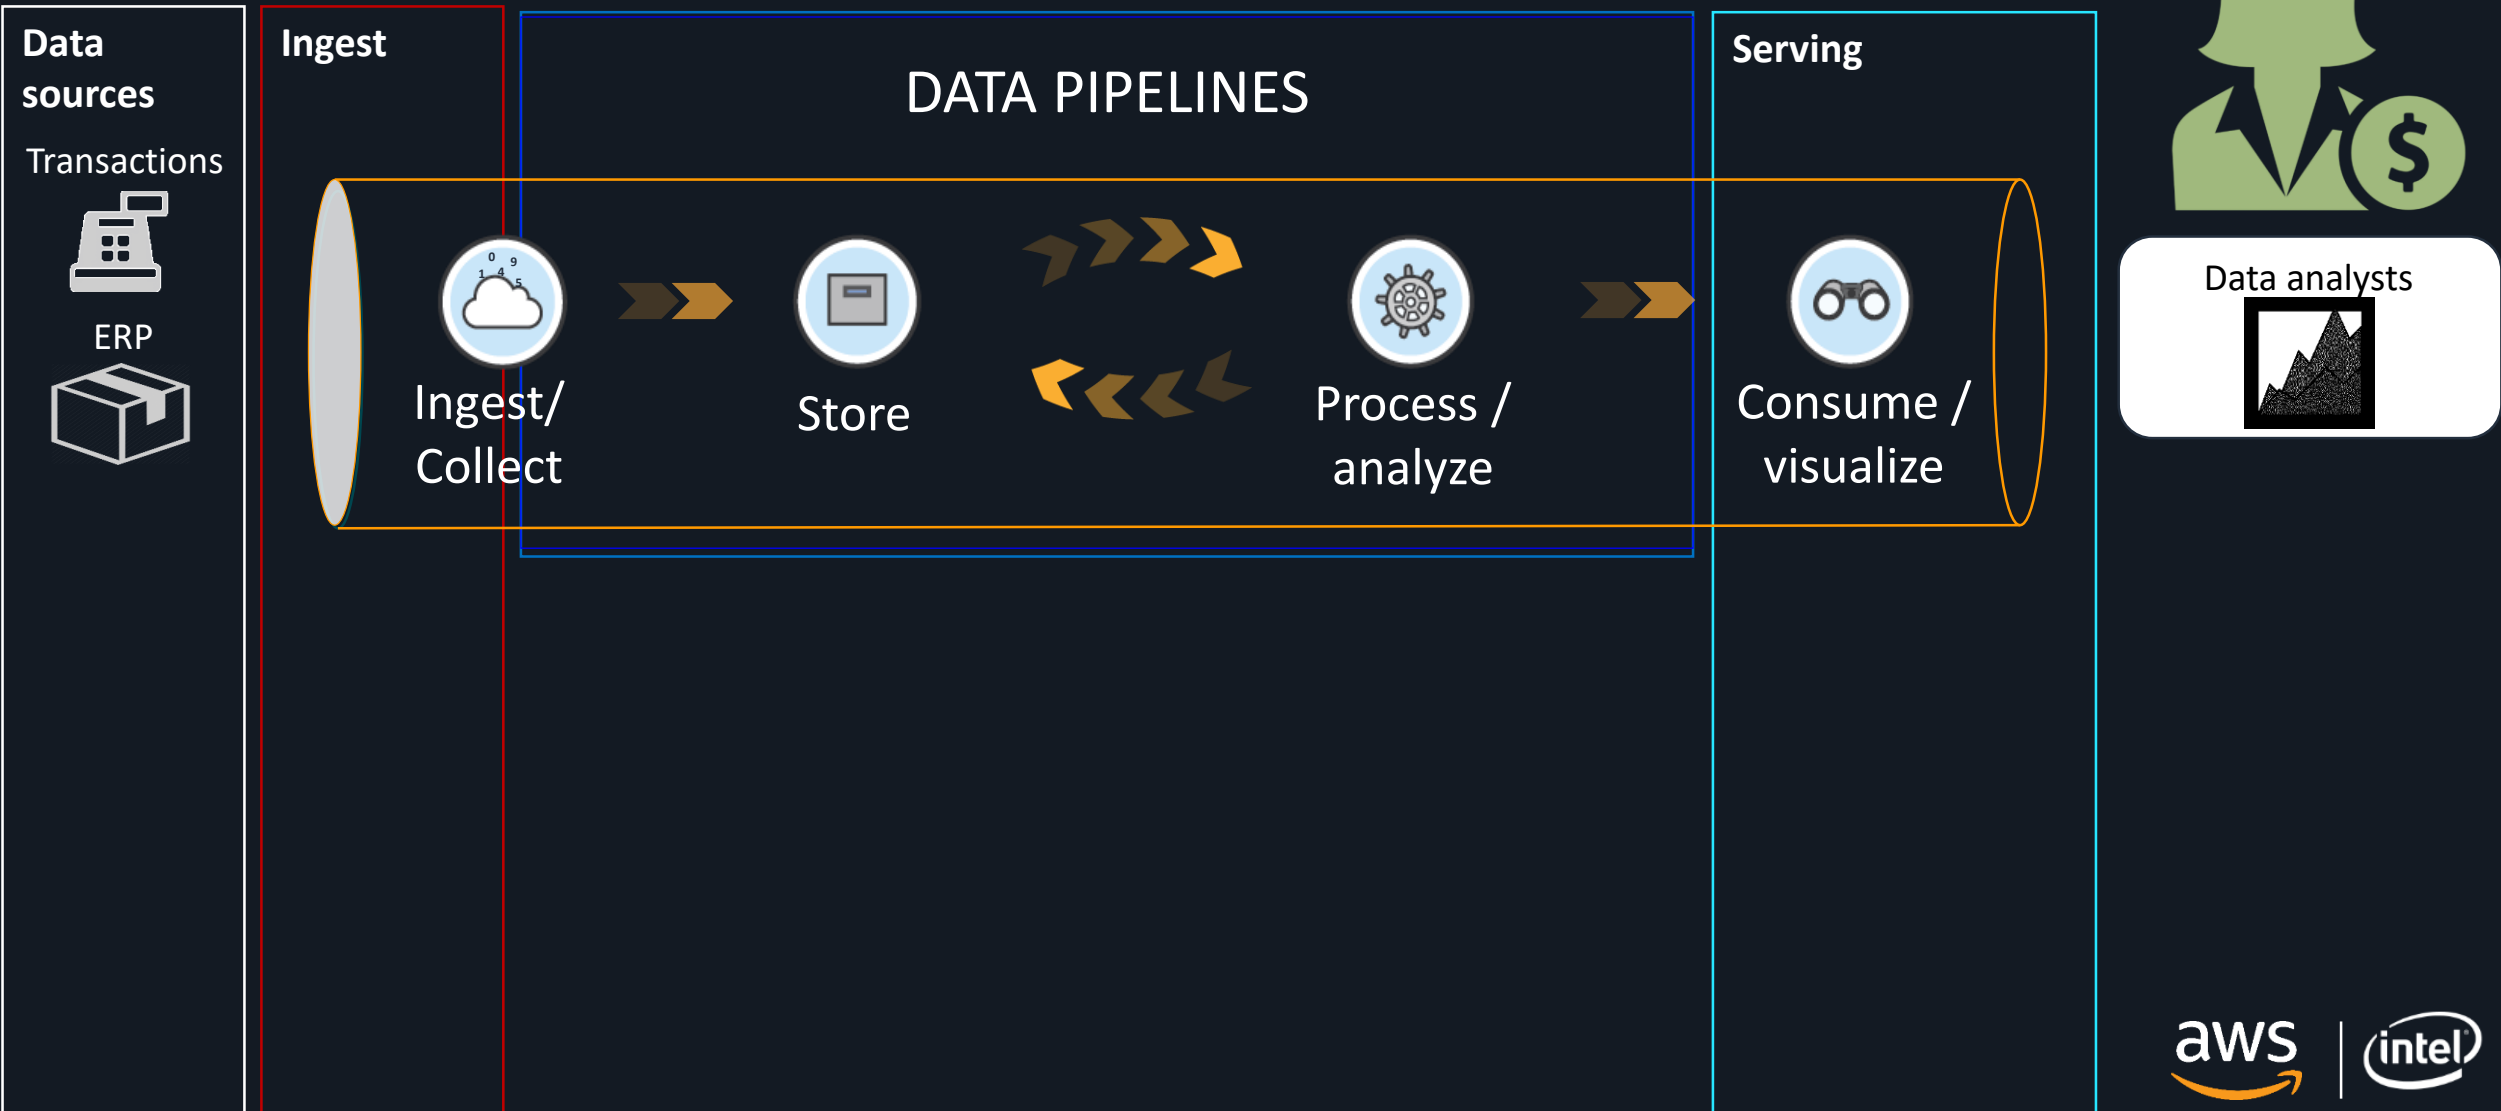
Order History /  
Returns (CRM)



Inventory /  
Production (ERP)

# Modern data architecture

Serving insights to users through data pipelines



# Start small and iterate

# Modern data architecture

Insights to enhance business applications, new digital services

Data  
sources

Transactions



ERP



SAP

Ingest

DATA PIPELINES

Serving



Data analysts

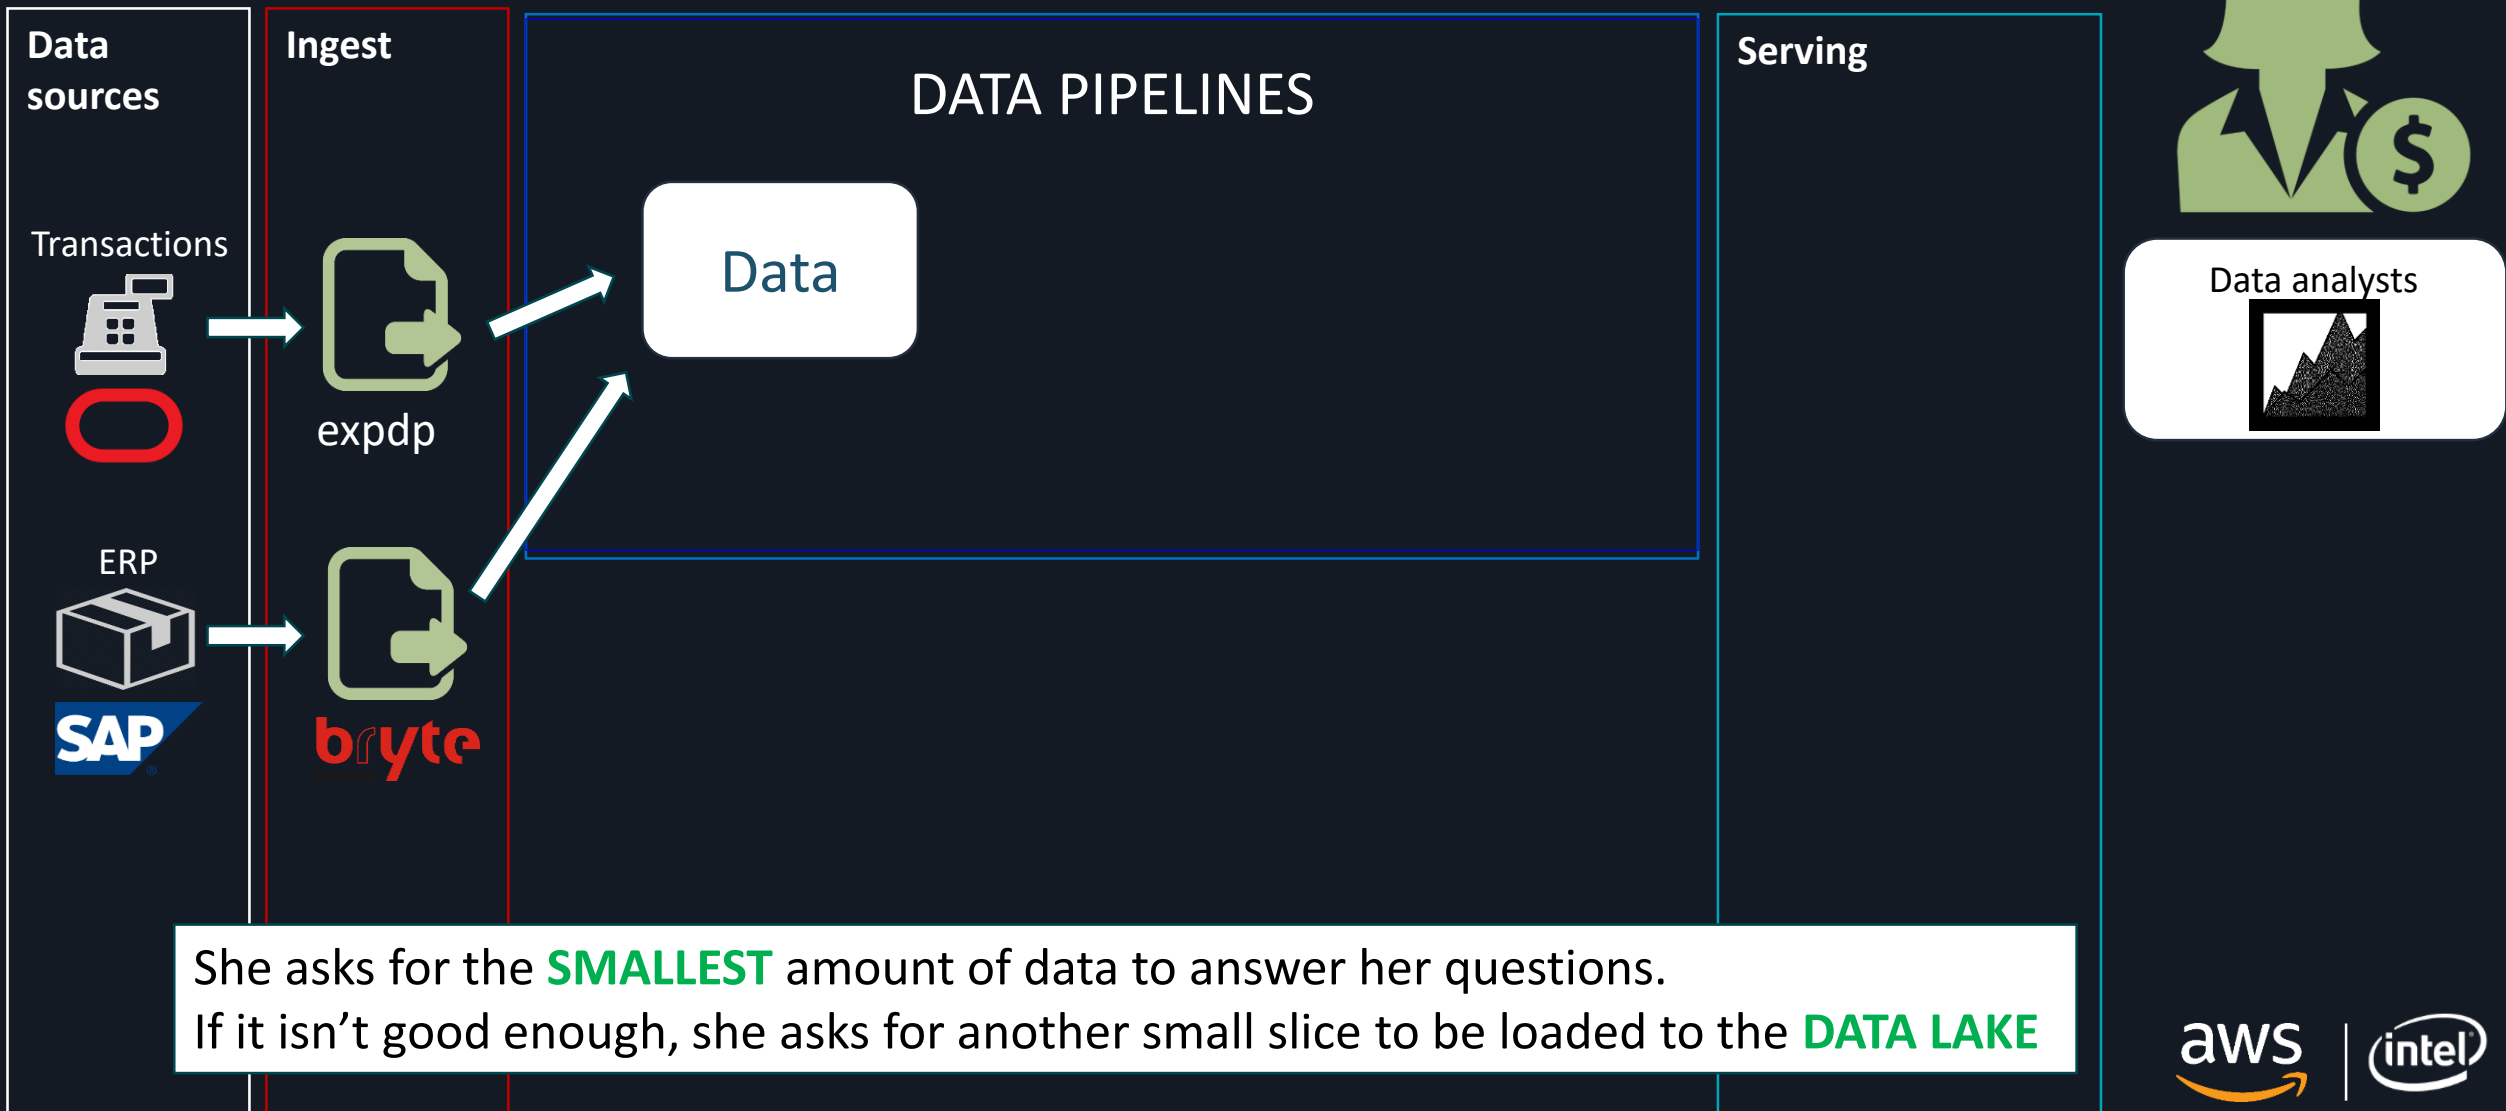


She asks for the **SMALLEST** amount of data to answer her questions.

If it isn't good enough, she asks for another small slice to be loaded to the **DATA LAKE**

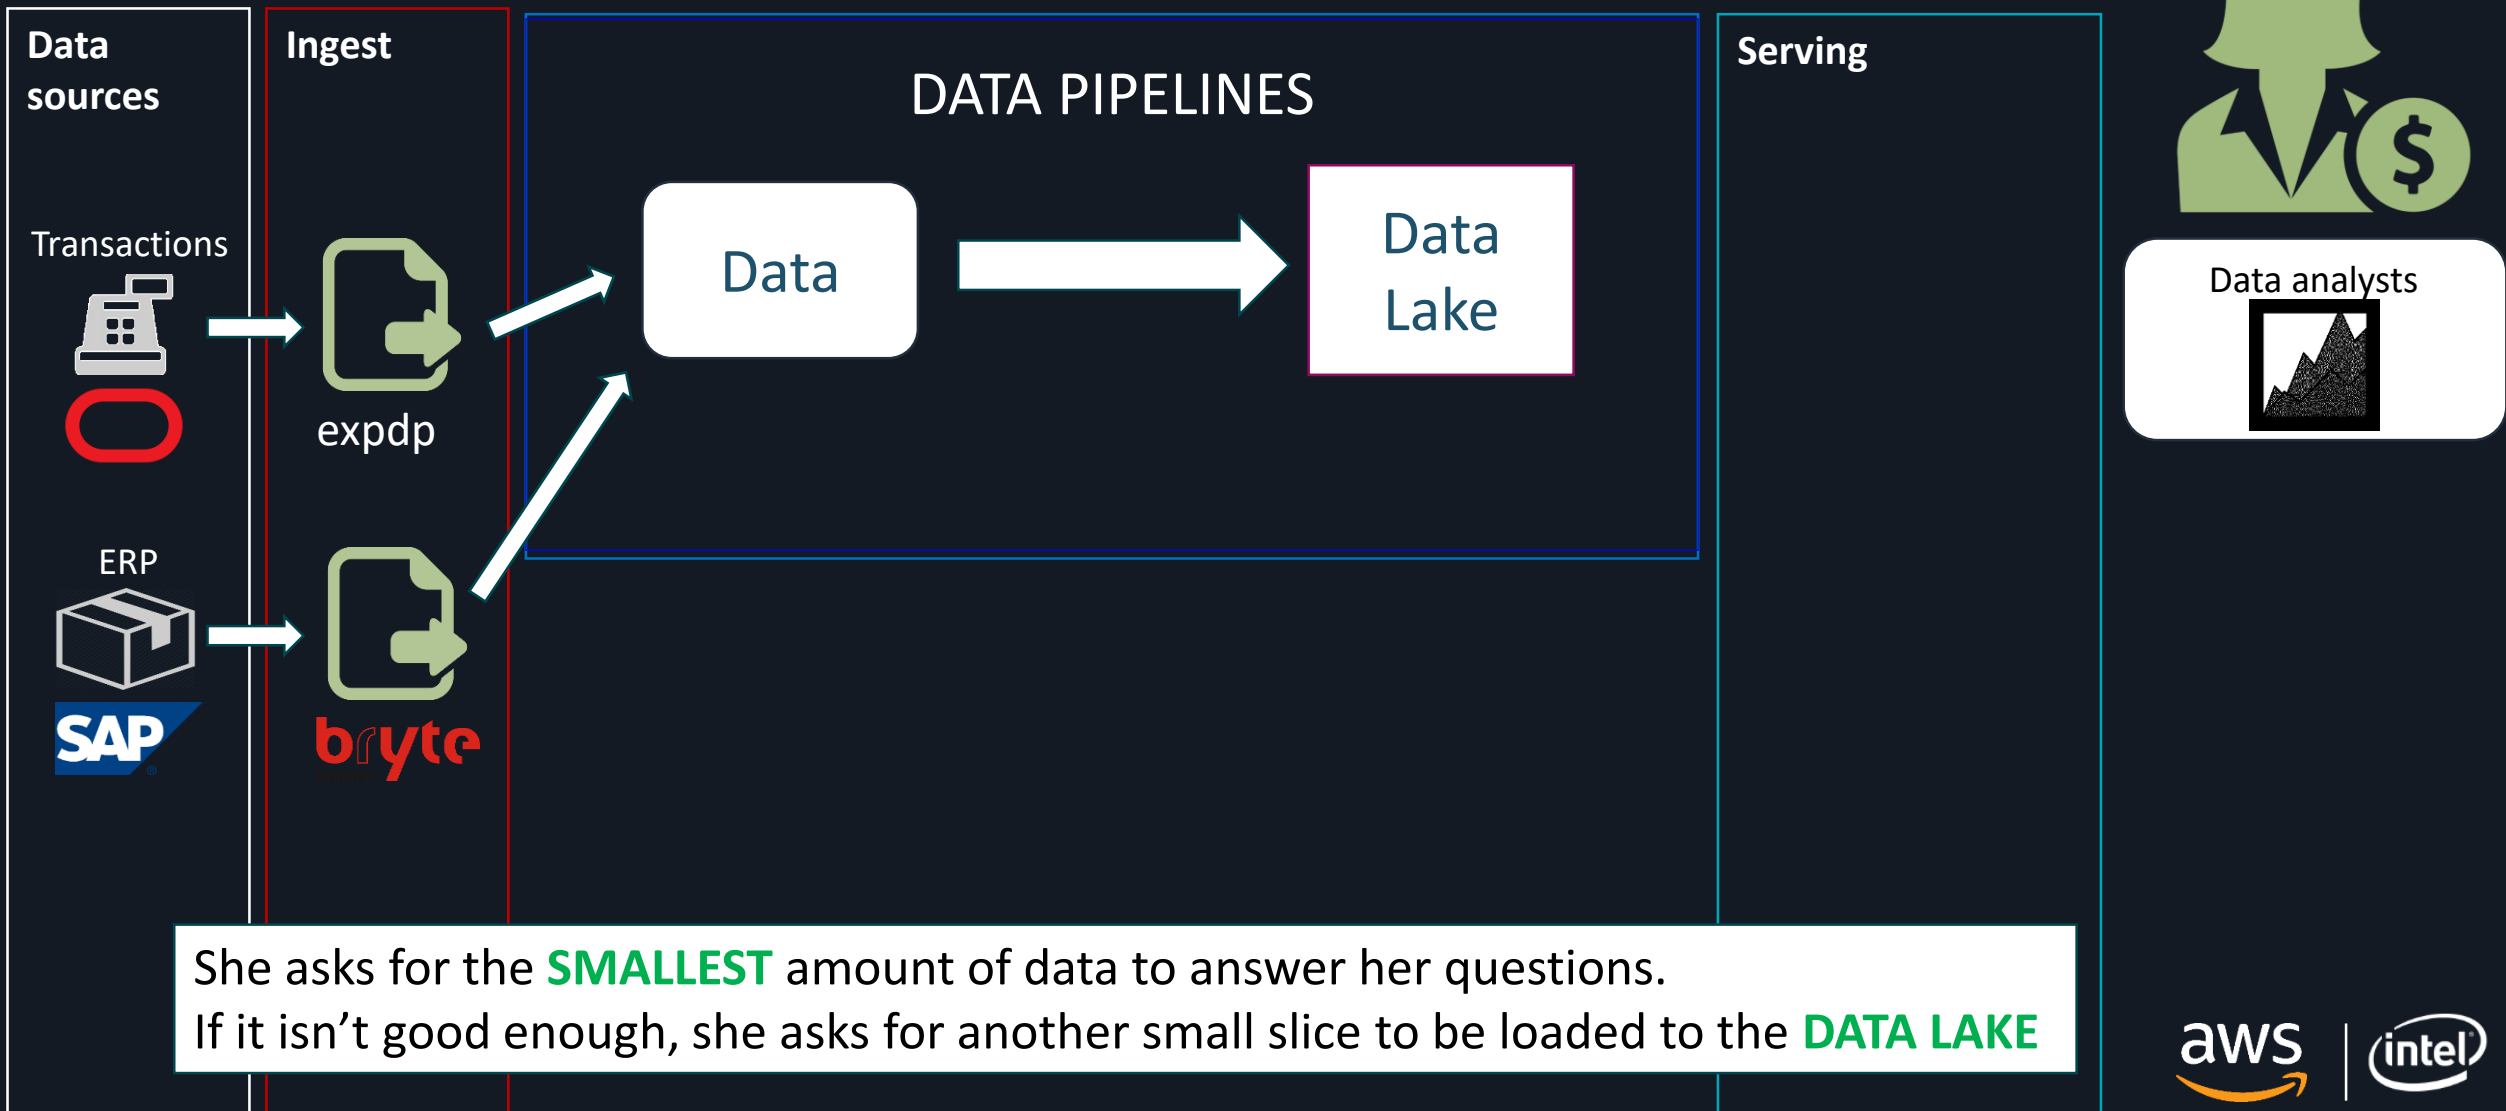
# Modern data architecture

Insights to enhance business applications, new digital services



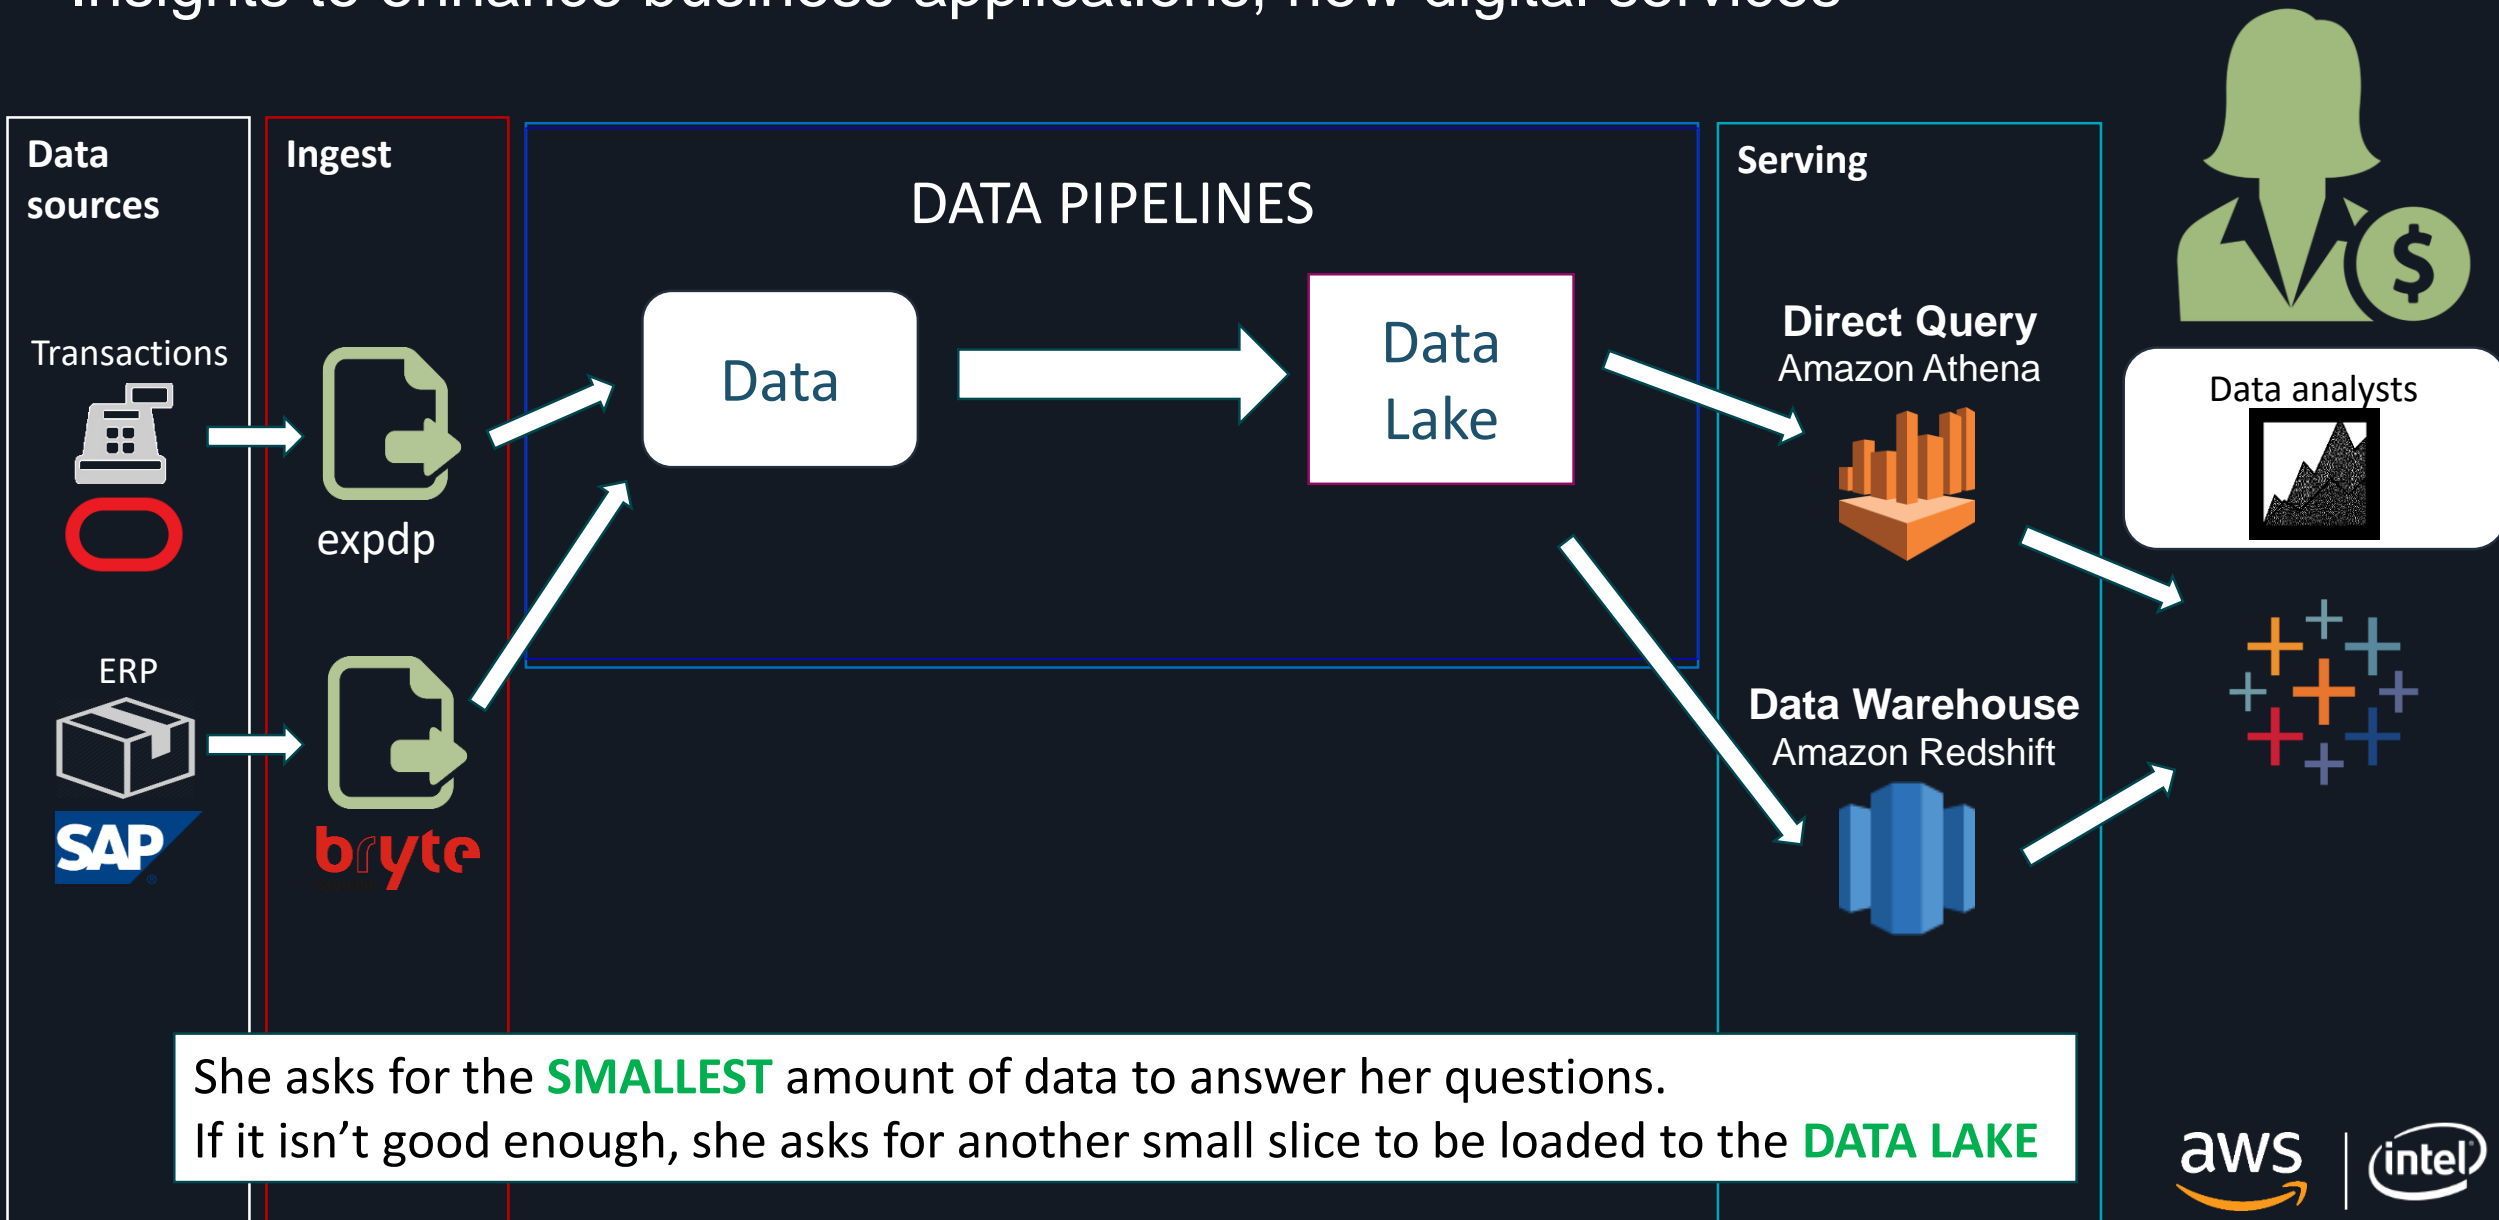
# Modern data architecture

Insights to enhance business applications, new digital services



# Modern data architecture

Insights to enhance business applications, new digital services





# Characteristics of a Data Lake



Collect  
Anything

# Characteristics of a Data Lake



Collect  
Anything



Dive in  
Anywhere


# Characteristics of a Data Lake



Collect  
Anything



Dive in  
Anywhere



Flexible  
Access


# Characteristics of a Data Lake



Collect  
Anything



Dive in  
Anywhere



Flexible  
Access



Future Proof

# Start with a set of specific questions to answer, then work backwards to the data required

Which regions are the new earbuds selling well?

What is the demand forecast by product category?

What is the social sentiment about our products?



Raymond, Head of Product

# Start with a set of specific questions to answer, then work backwards to the data required

Which **regions** are the new earbuds selling well?

What is the **demand forecast** by product category?

What is the **social sentiment** about our products?



Raymond, Head of Product



Order History /  
Returns (CRM)

**NOW IN THE DATA LAKE**

# Start with a set of specific questions to answer, then work backwards to the data required

Which regions are the new earbuds selling well?

What is the demand forecast by product category?

What is the social sentiment about our products?



Raymond, Head of Product



Order History /  
Returns (CRM)

NOW IN THE DATA LAKE



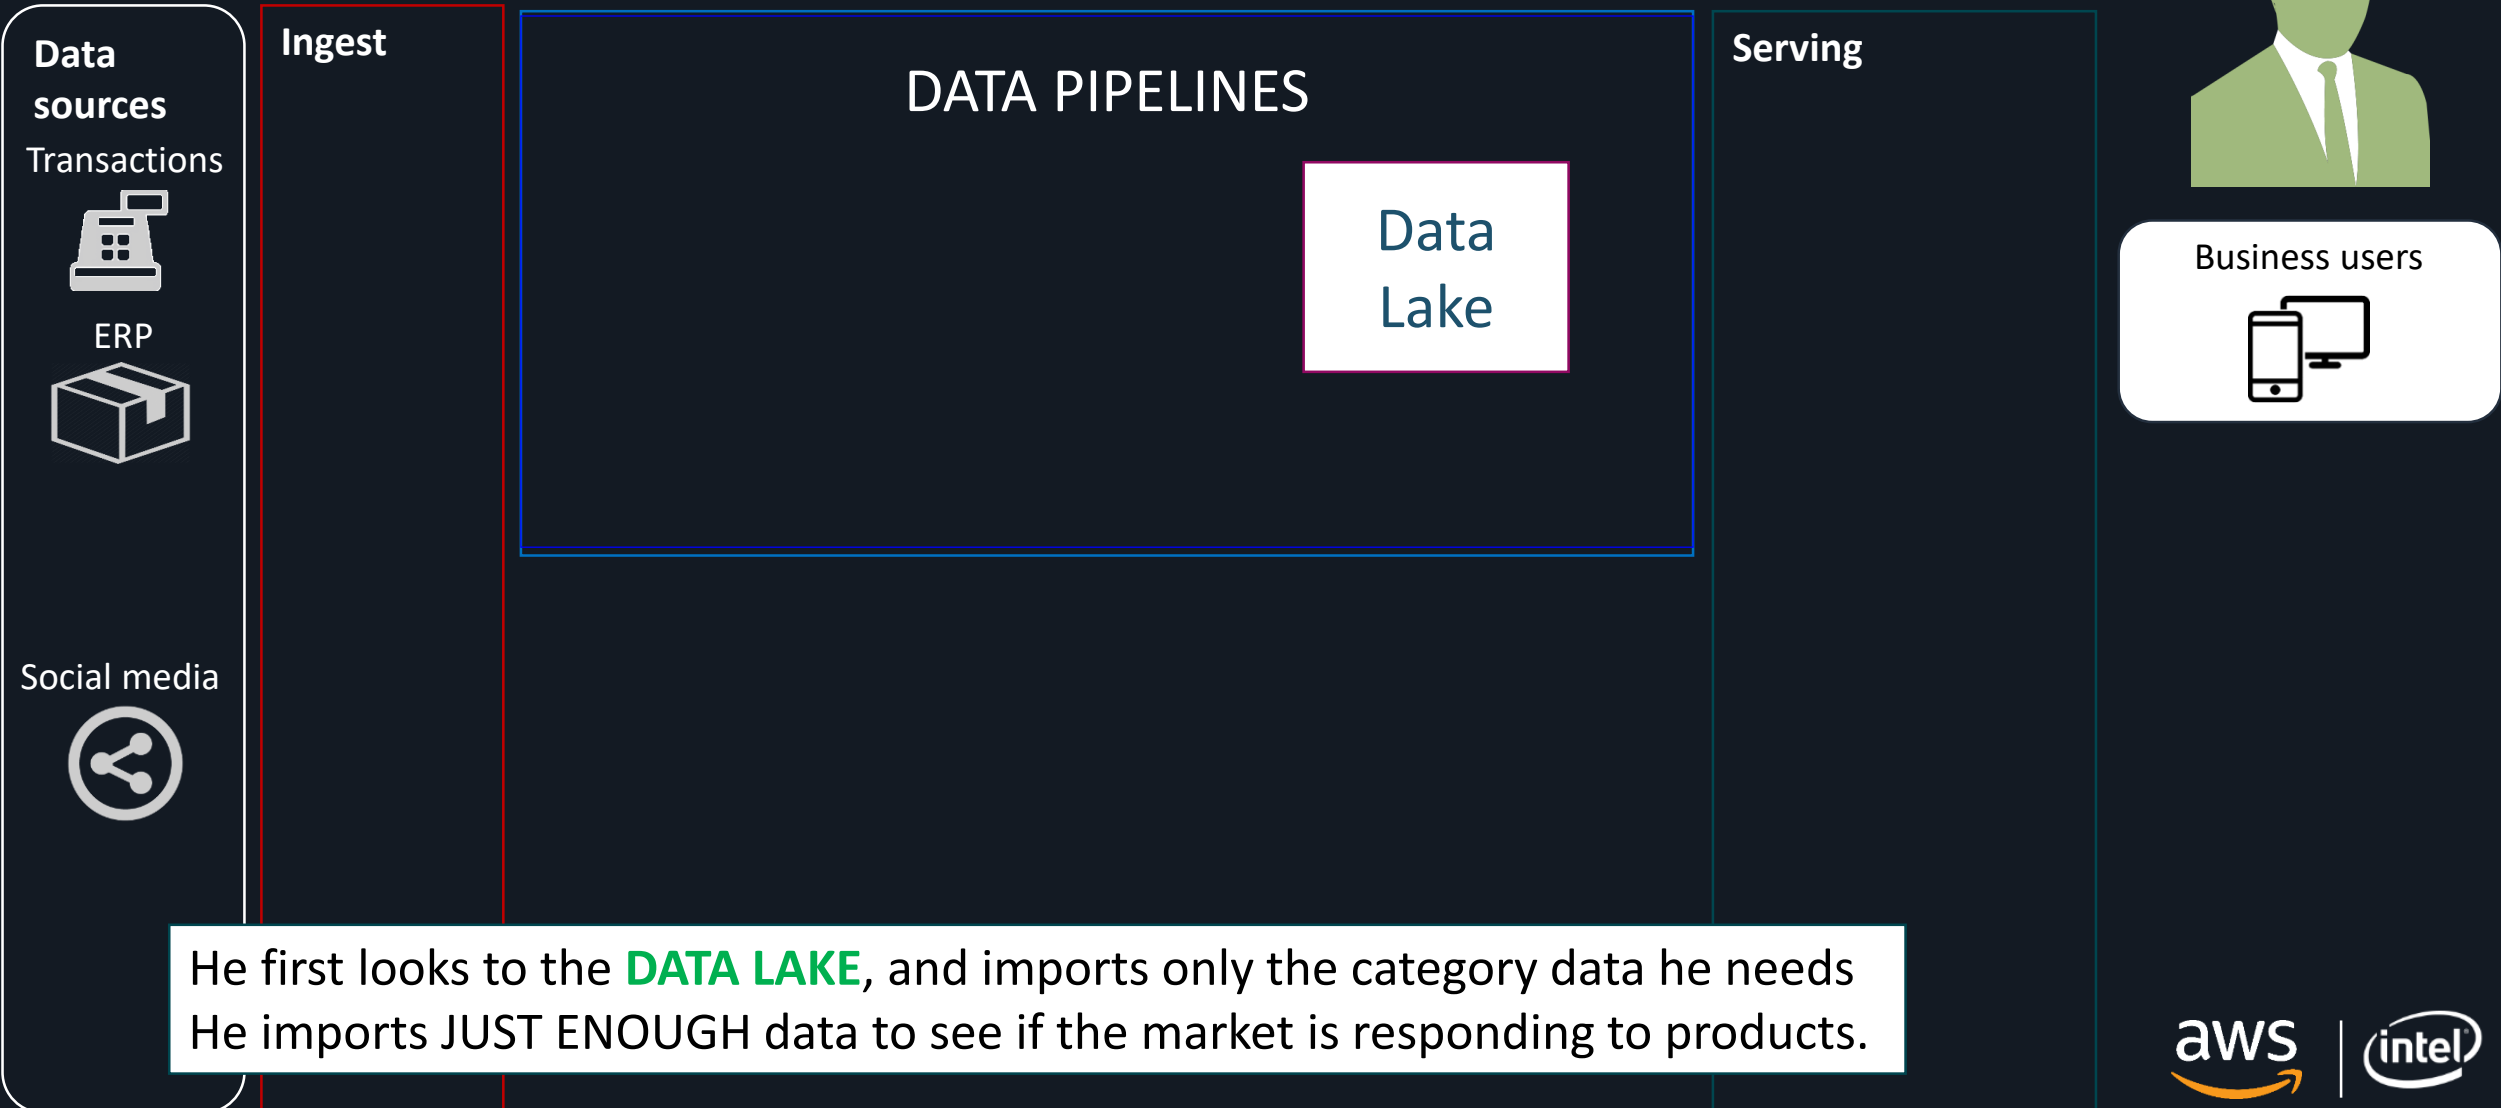
Trending /  
Mentions (Social)

# Experiment, validate, then scale



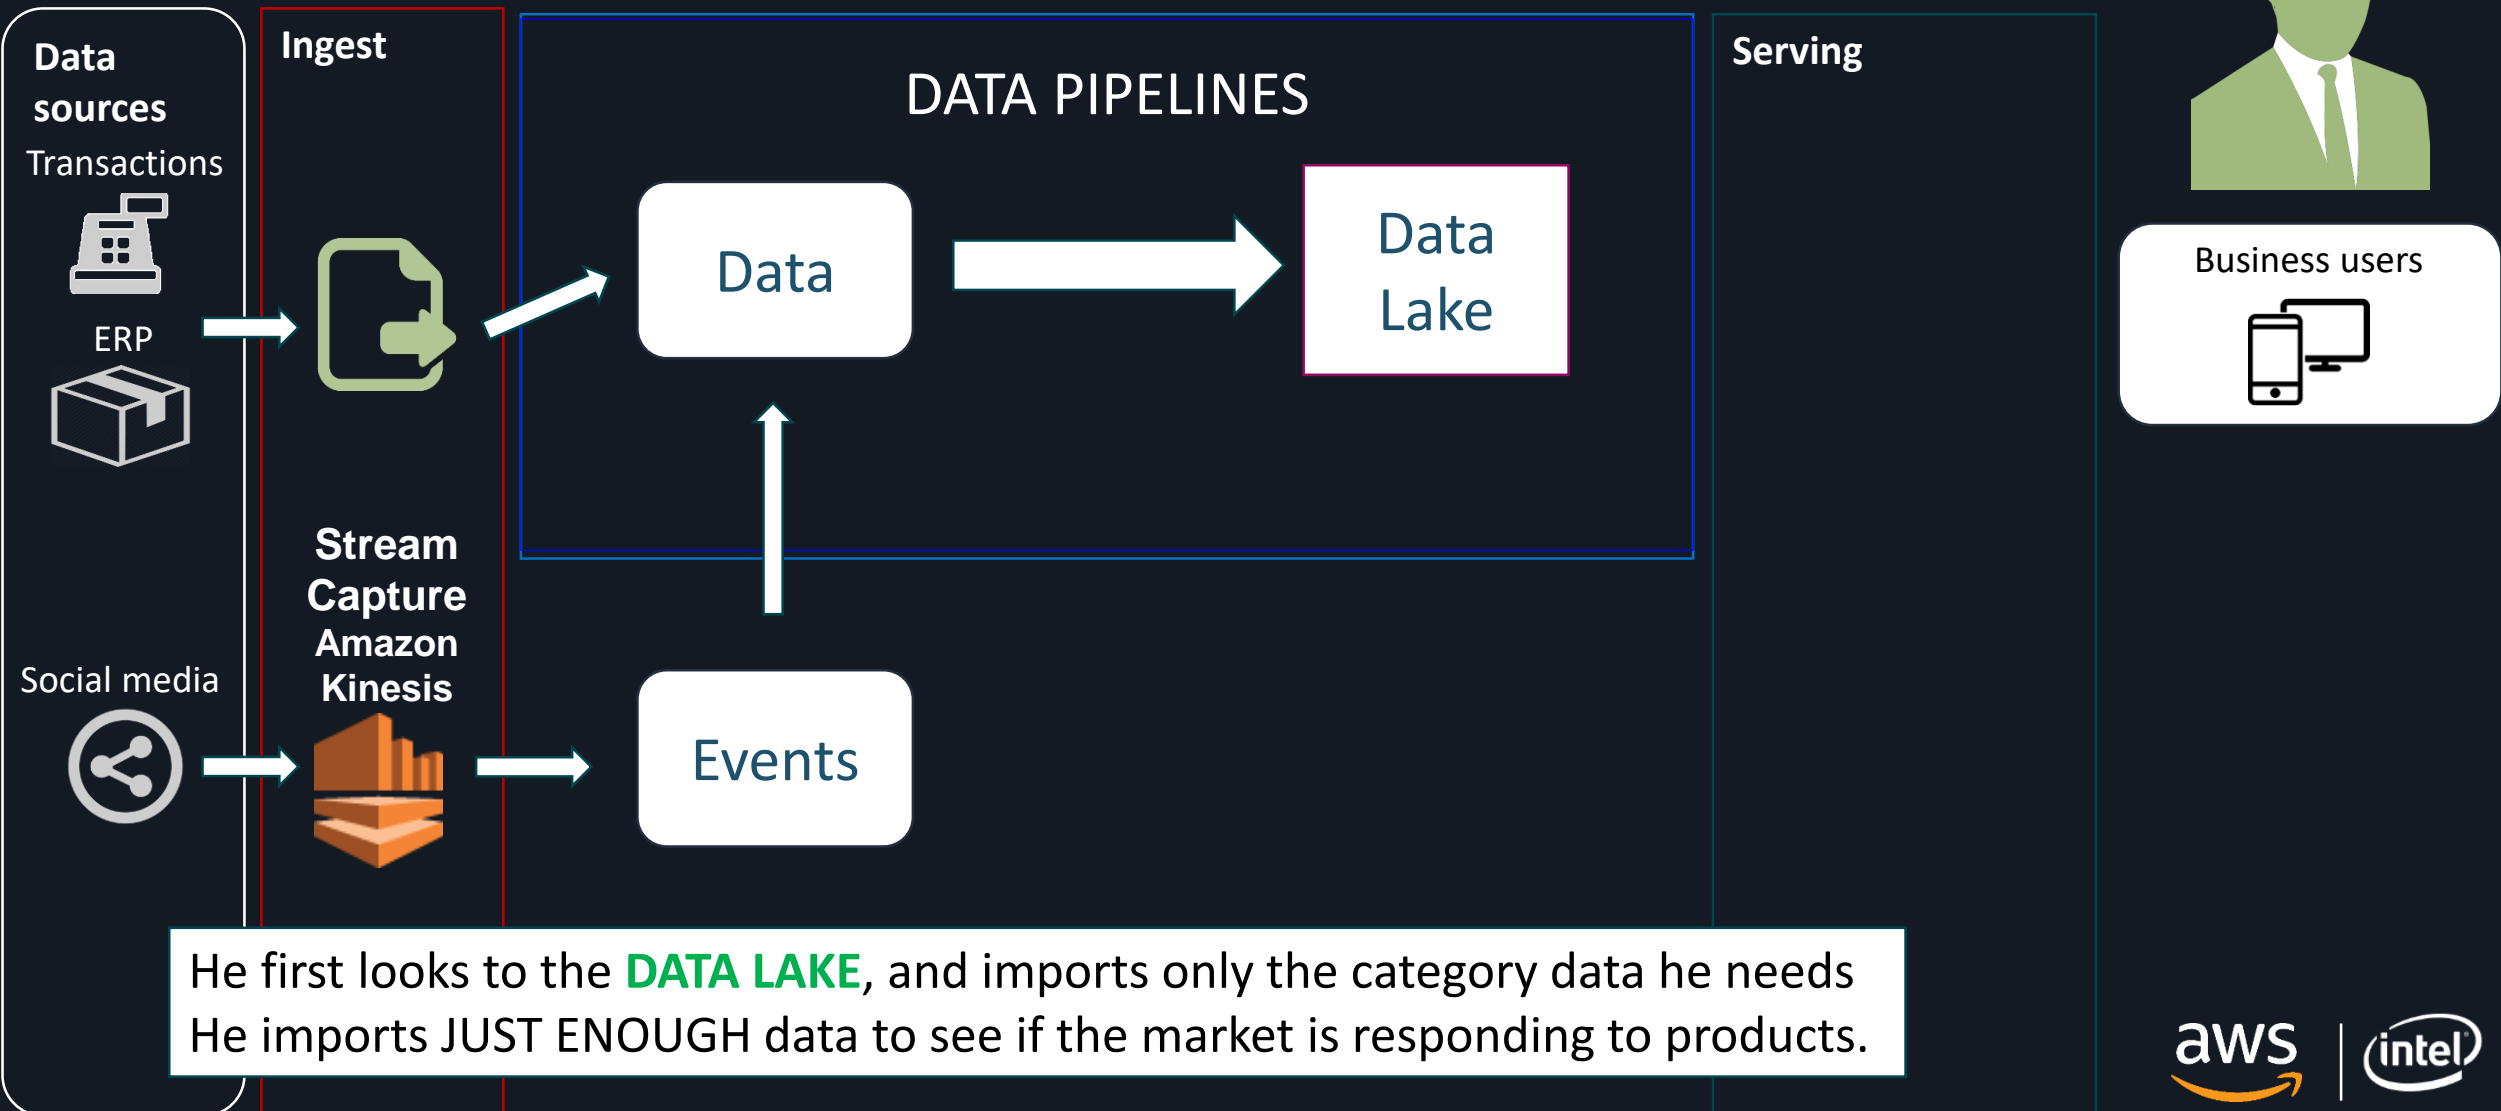
# Modern data architecture

Insights to enhance business applications, new digital services



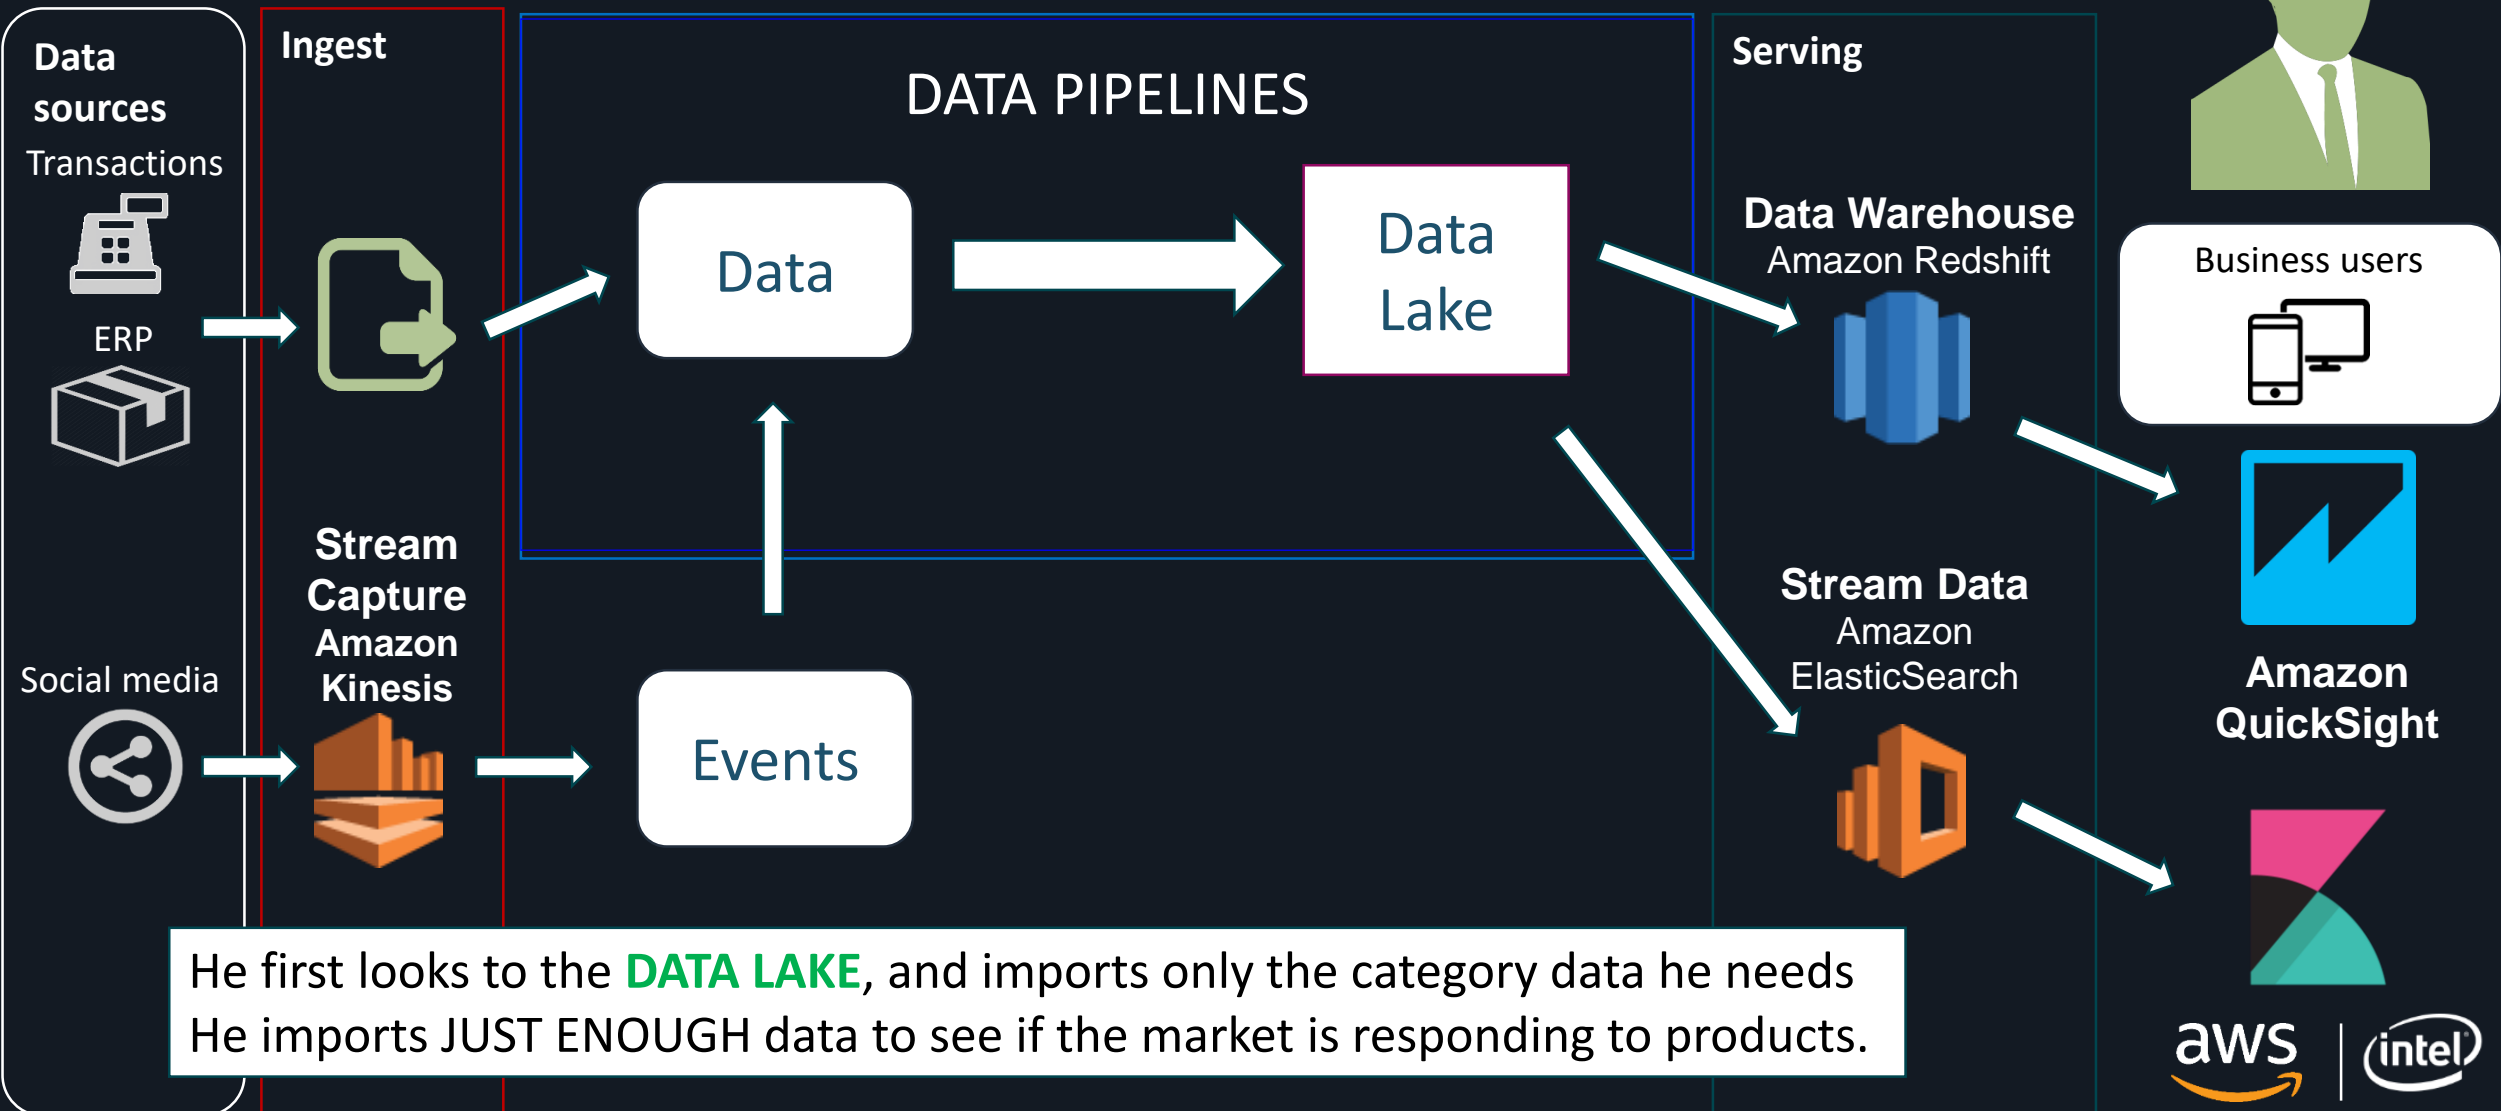
# Modern data architecture

Insights to enhance business applications, new digital services



# Modern data architecture

Insights to enhance business applications, new digital services



# Amazon Redshift – Modern Data Warehousing

Fast, scalable, fully managed data warehouse at 1/10<sup>th</sup> the cost

Massively parallel, scales from gigabytes to exabytes

Queries data across your Redshift data warehouse and Amazon S3 data lake

---

## Fast at scale



Columnar storage technology to improve I/O efficiency and scale query performance

## Open file formats



Analyze optimized data formats on direct-attached disks, and all open file formats in S3

## Secure



Audit everything; encrypt data end-to-end; extensive certification and compliance

## Cost-effective



Start at \$0.25 per hour; as low as \$250-\$333 per uncompressed terabyte per year


# Common data pipeline configuration

Highly decoupled configurations scale better, are more fault tolerant, and cost optimized

# Common data pipeline configuration


Highly decoupled configurations scale better, are more fault tolerant, and cost optimized

Raw Data  
Amazon S3



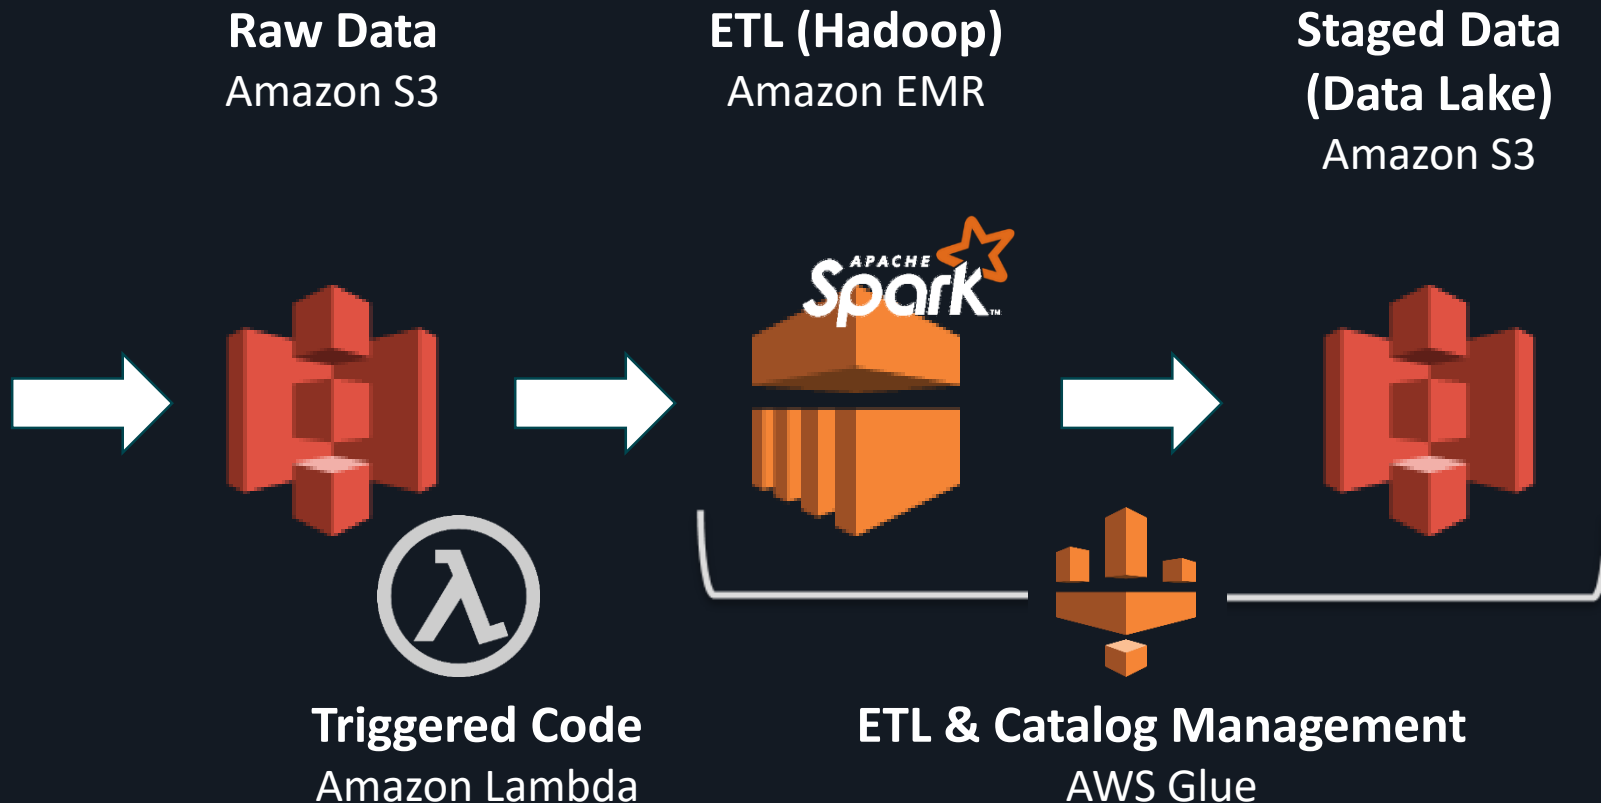
# Common data pipeline configuration

Highly decoupled configurations scale better, are more fault tolerant, and cost optimized



# Common data pipeline configuration

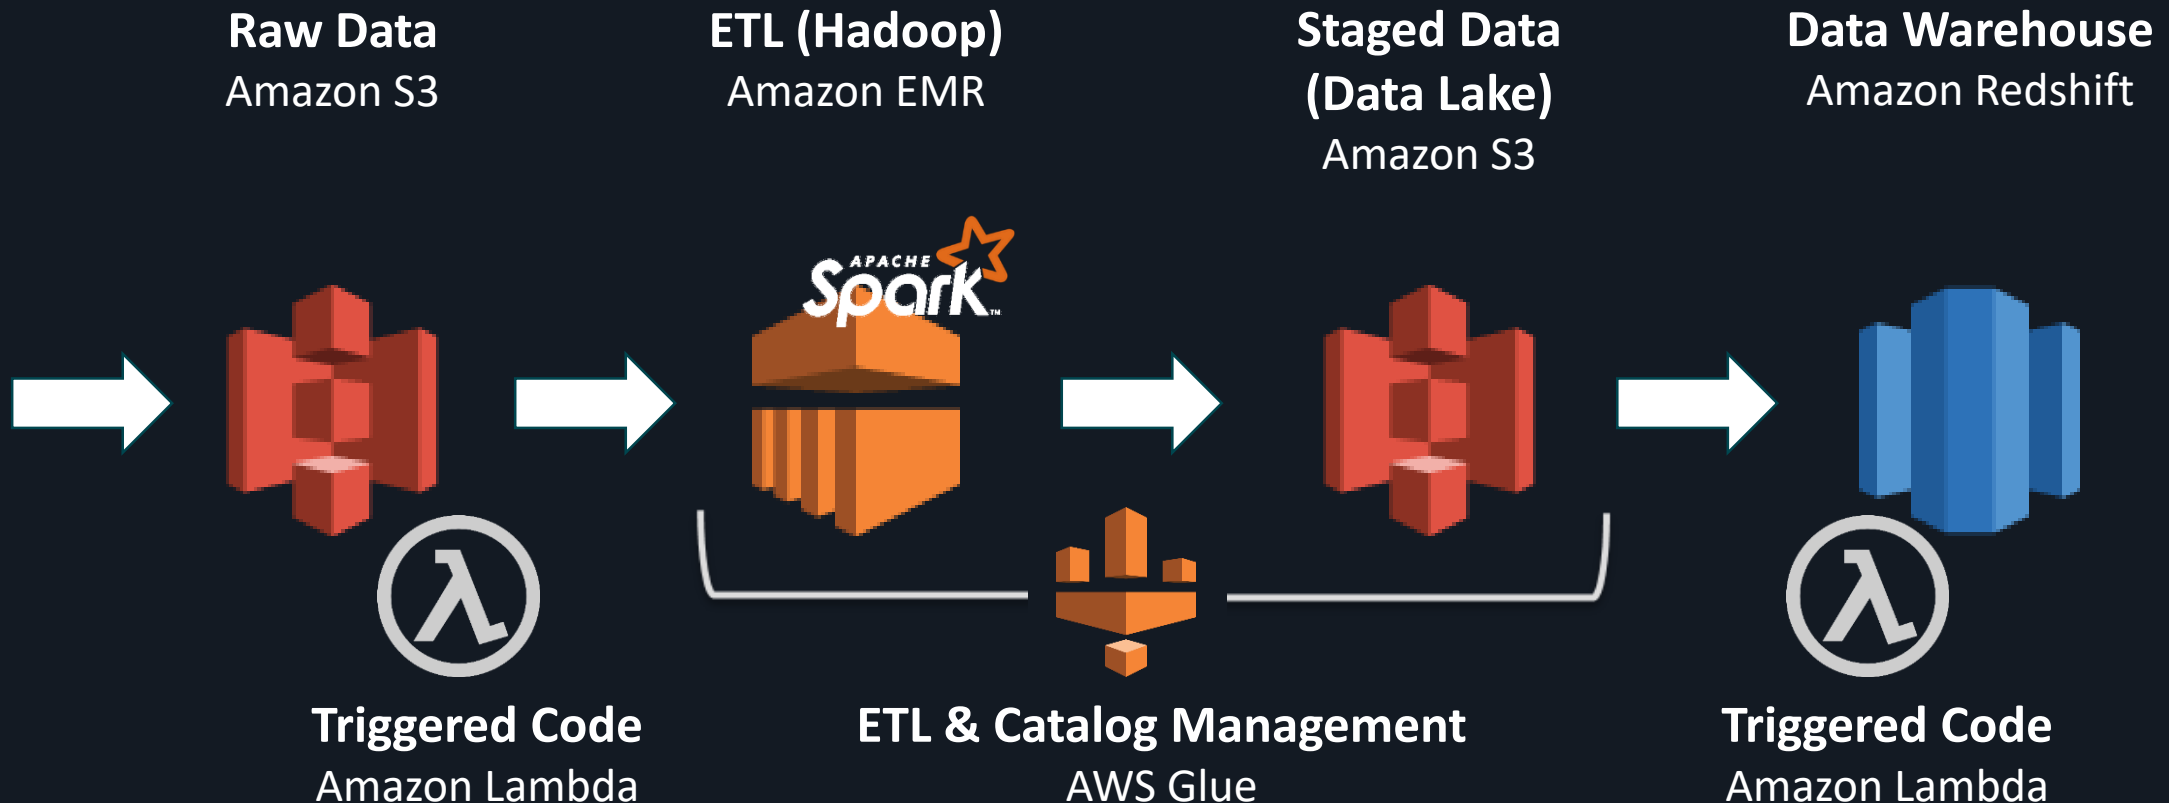
Highly decoupled configurations scale better, are more fault tolerant, and cost optimized





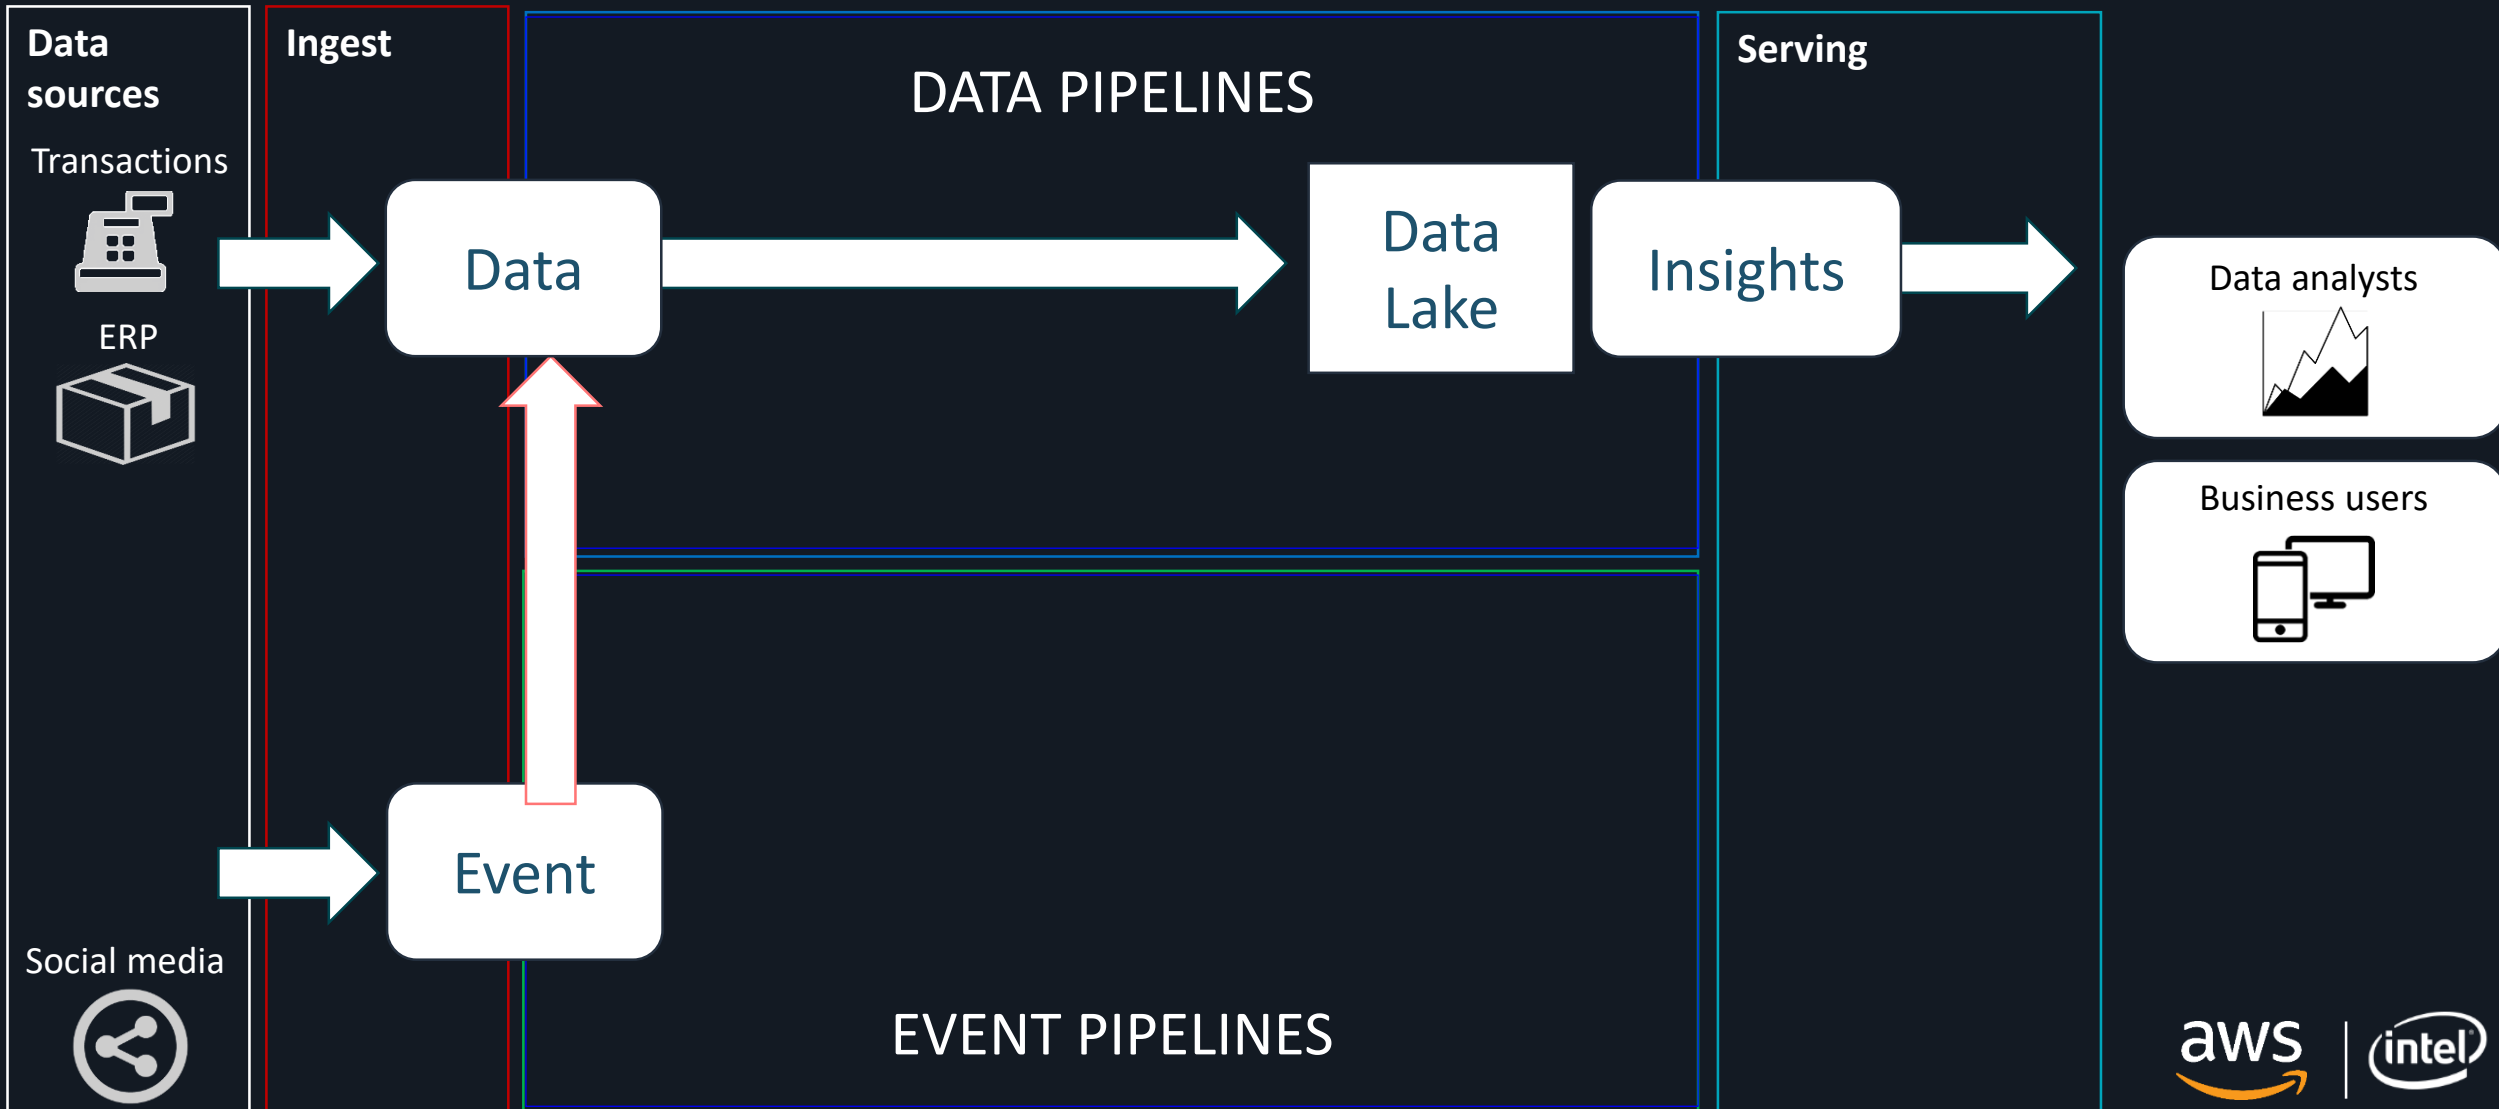
# Common data pipeline configuration

Highly decoupled configurations scale better, are more fault tolerant, and cost optimized



# Modern data architecture

Insights to enhance business applications, new digital services



# Outcome 2: Innovate for new revenues

# EarEcstasy has its first direct relationship with consumers



Bala, Head of Marketing



Krzysztof, Data Scientist

# EarEcstasy has its first direct relationship with consumers

How are people using the Smart Buds?

How to understand what they listen to and when?

What kinds of people are in/decreasing usage?



Bala, Head of Marketing



Krzysztof, Data Scientist

# EarEcstasy has its first direct relationship with consumers

How are people using the Smart Buds?

How to understand what they listen to and when?

What kinds of people are in/decreasing usage?



Bala, Head of Marketing

What are our customer segments, based on usage?

Can predict music preference from location and HR?

Are there additional signals in the voice commands?

Can we infer user activity, from scenes in pictures?



Krzysztof, Data Scientist

# Start with a set of specific questions to answer, then work backwards to the data required

How are people using the Smart Buds?

How to understand what they listen to and when?

What kinds of people are in/decreasing usage?



Bala, Head of Marketing

# Start with a set of specific questions to answer, then work backwards to the data required

How are people using the Smart Buds?

How to understand what they listen to and when?

What kinds of people are in/decreasing usage?



Bala, Head of Marketing



Registration, usage  
[time/place] (Mobile  
app)



# Start with a set of specific questions to answer, then work backwards to the data required

How are people using the Smart Buds?

How to understand what they listen to and when?

What kinds of people are in/decreasing usage?



Bala, Head of Marketing



Registration, usage  
[time/place] (Mobile  
app)



Media consumption  
(Partner API)



# Start with a set of specific questions to answer, then work backwards to the data required

What are our **customer segments**, based on usage?

Can **predict music preference** from location and HR?

Are there additional **signals in the voice** commands?

Can we **infer user activity**, from scenes in pictures?



Krzysztof, Data Scientist

# Start with a set of specific questions to answer, then work backwards to the data required

What are our **customer segments**, based on usage?

Can **predict music preference** from location and HR?

Are there additional **signals in the voice** commands?

Can we **infer user activity**, from scenes in pictures?



Krzysztof, Data Scientist



# Start with a set of specific questions to answer, then work backwards to the data required

What are our **customer segments**, based on usage?

Can **predict music preference** from location and HR?

Are there additional **signals in the voice** commands?

Can we **infer user activity**, from scenes in pictures?



Krzysztof, Data Scientist



Registration, usage  
[time/place] (Mobile  
app)

**LOAD TO DATA LAKE**



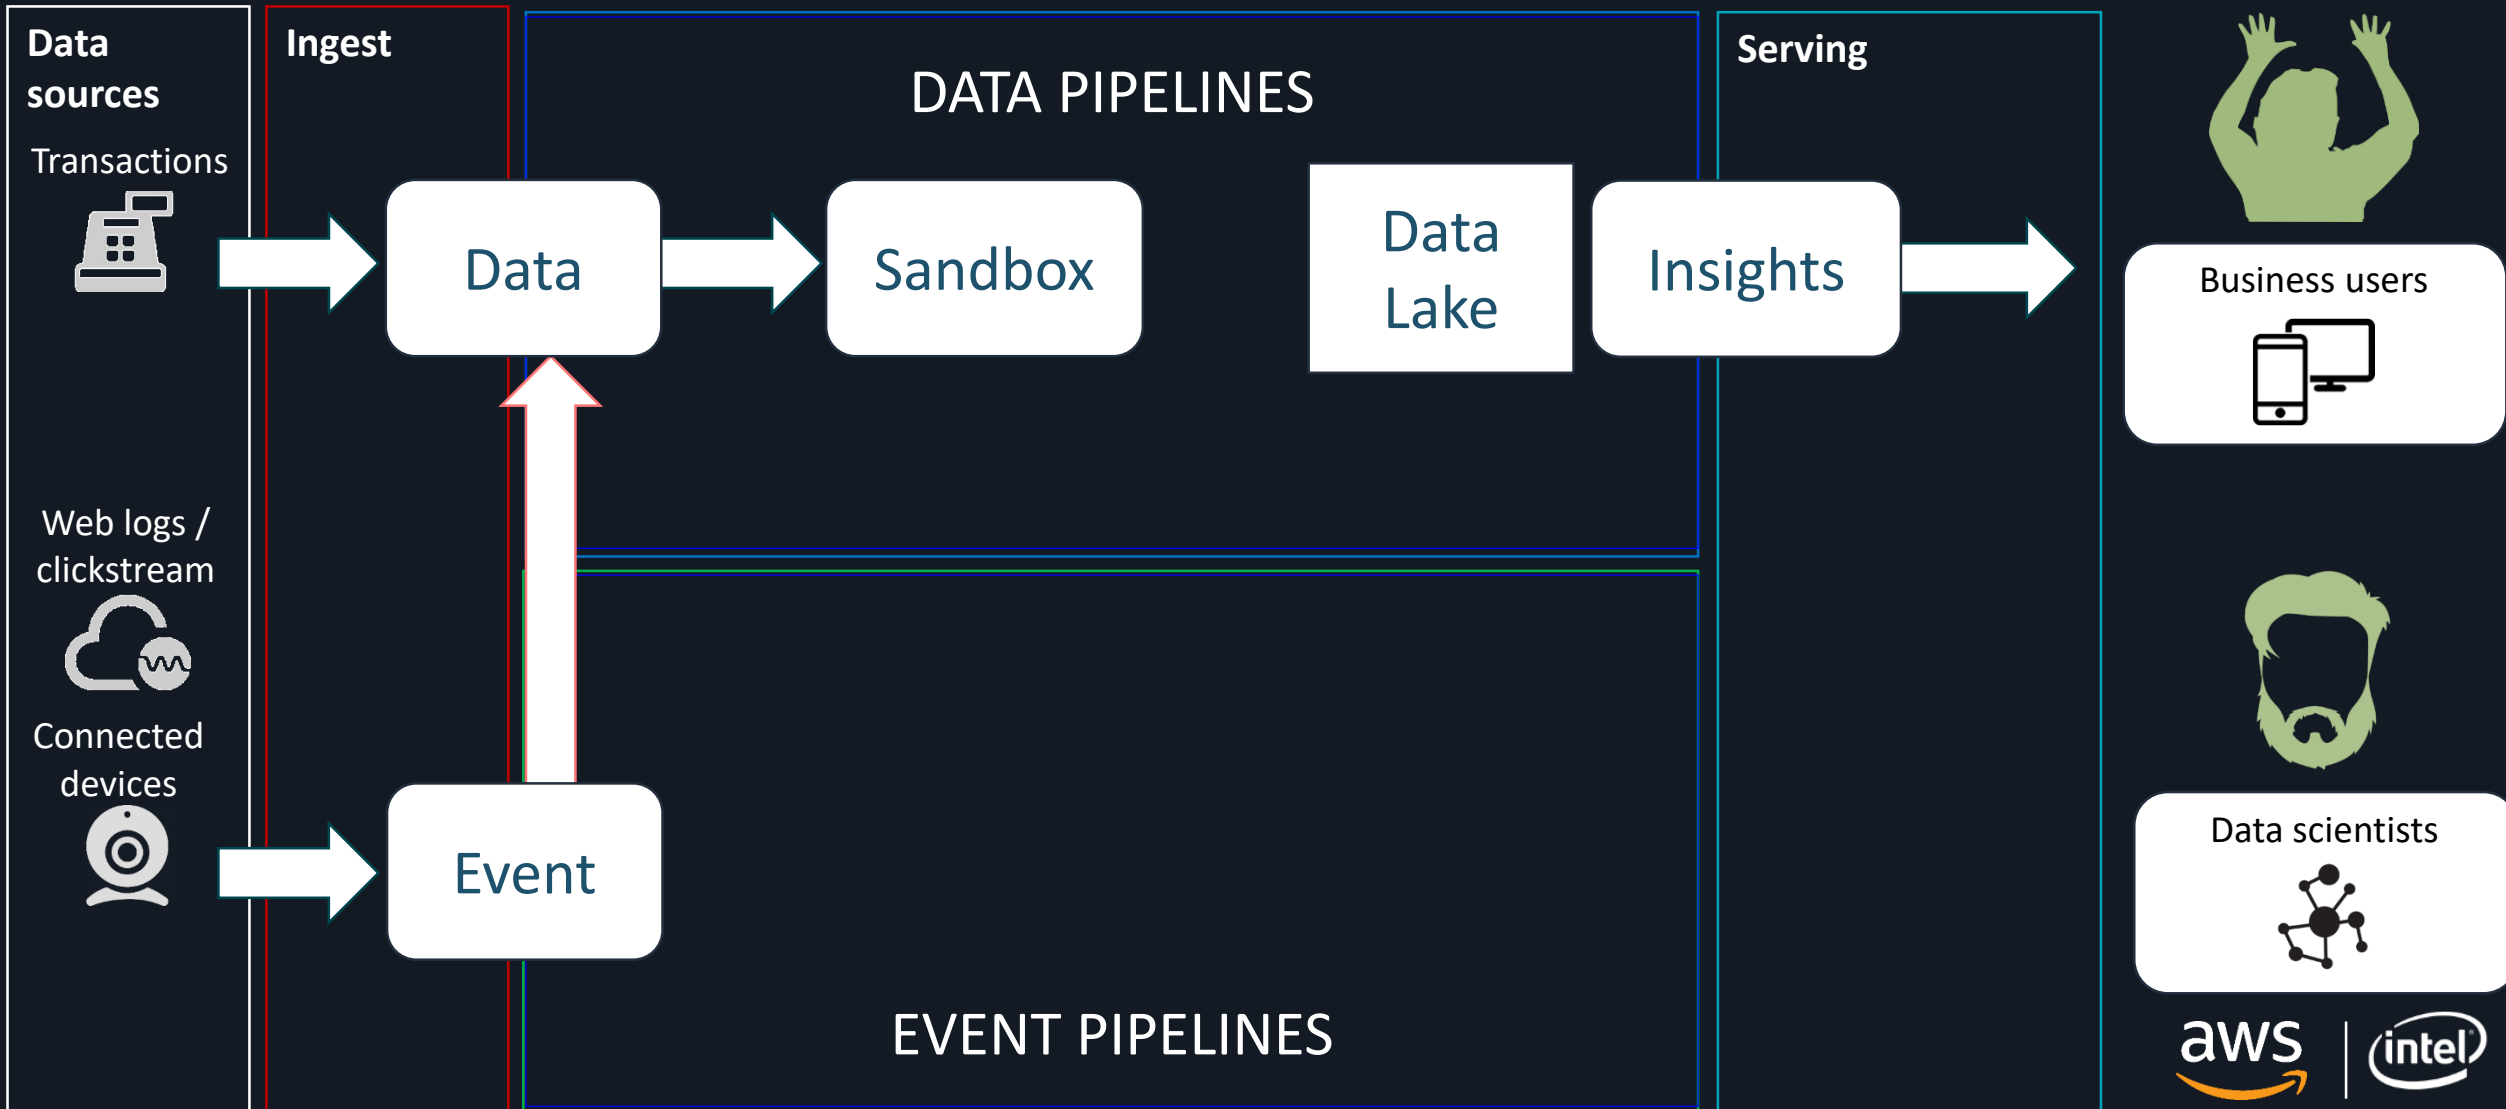
HR, Voice, GPS,  
Images (Device  
data)

**DATA LAKE, OR NOT?**

# Sandboxes - fast, cheap, low risk

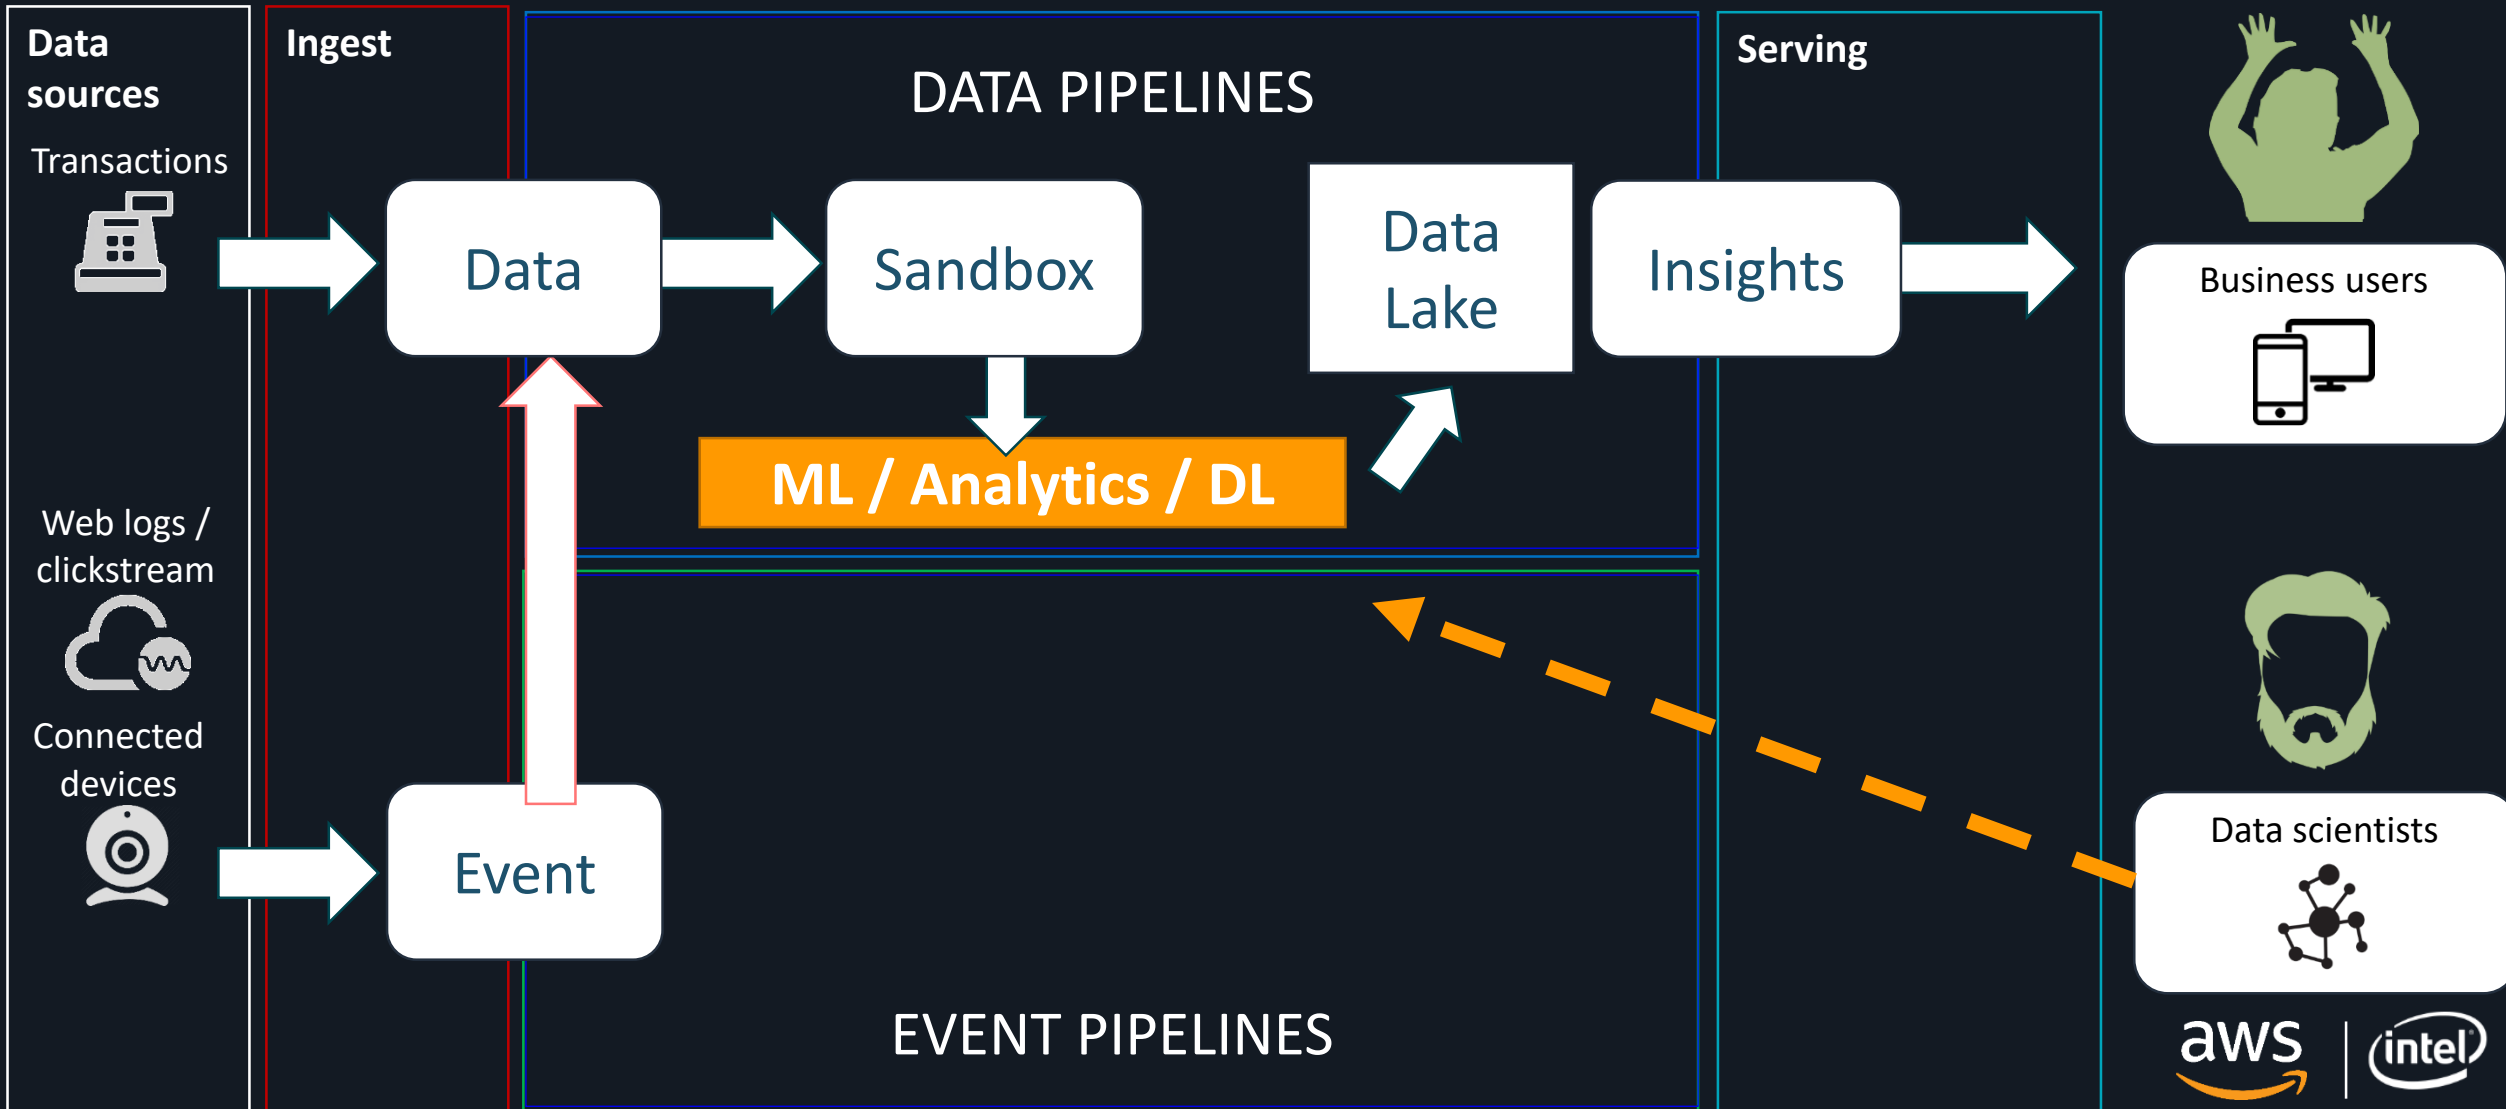
# Modern data architecture

Insights to enhance business applications, new digital services



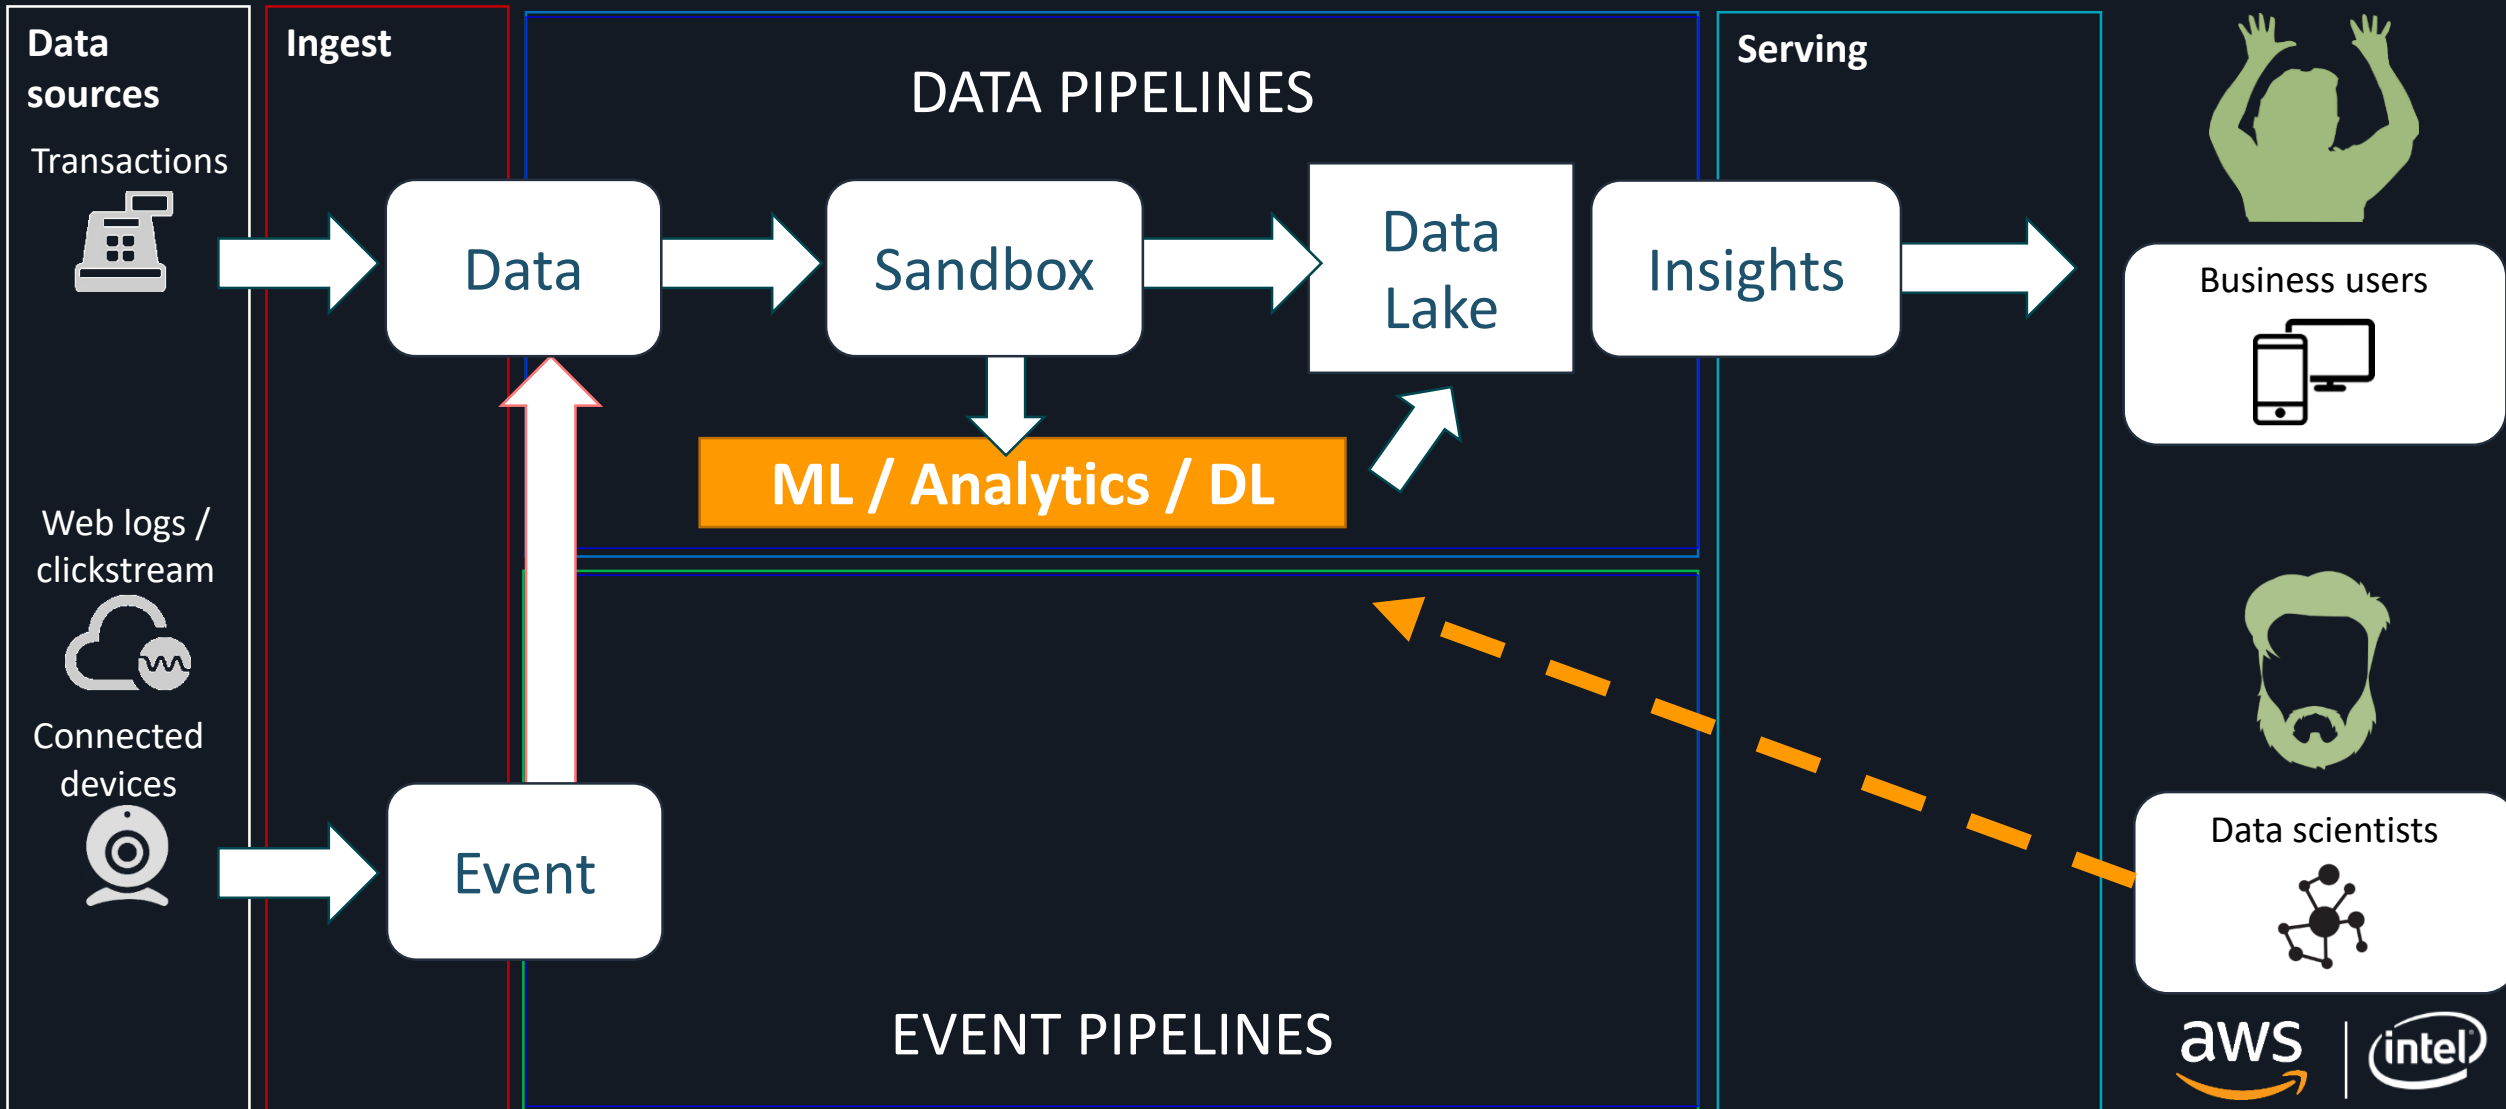
# Modern data architecture

Insights to enhance business applications, new digital services



# Modern data architecture

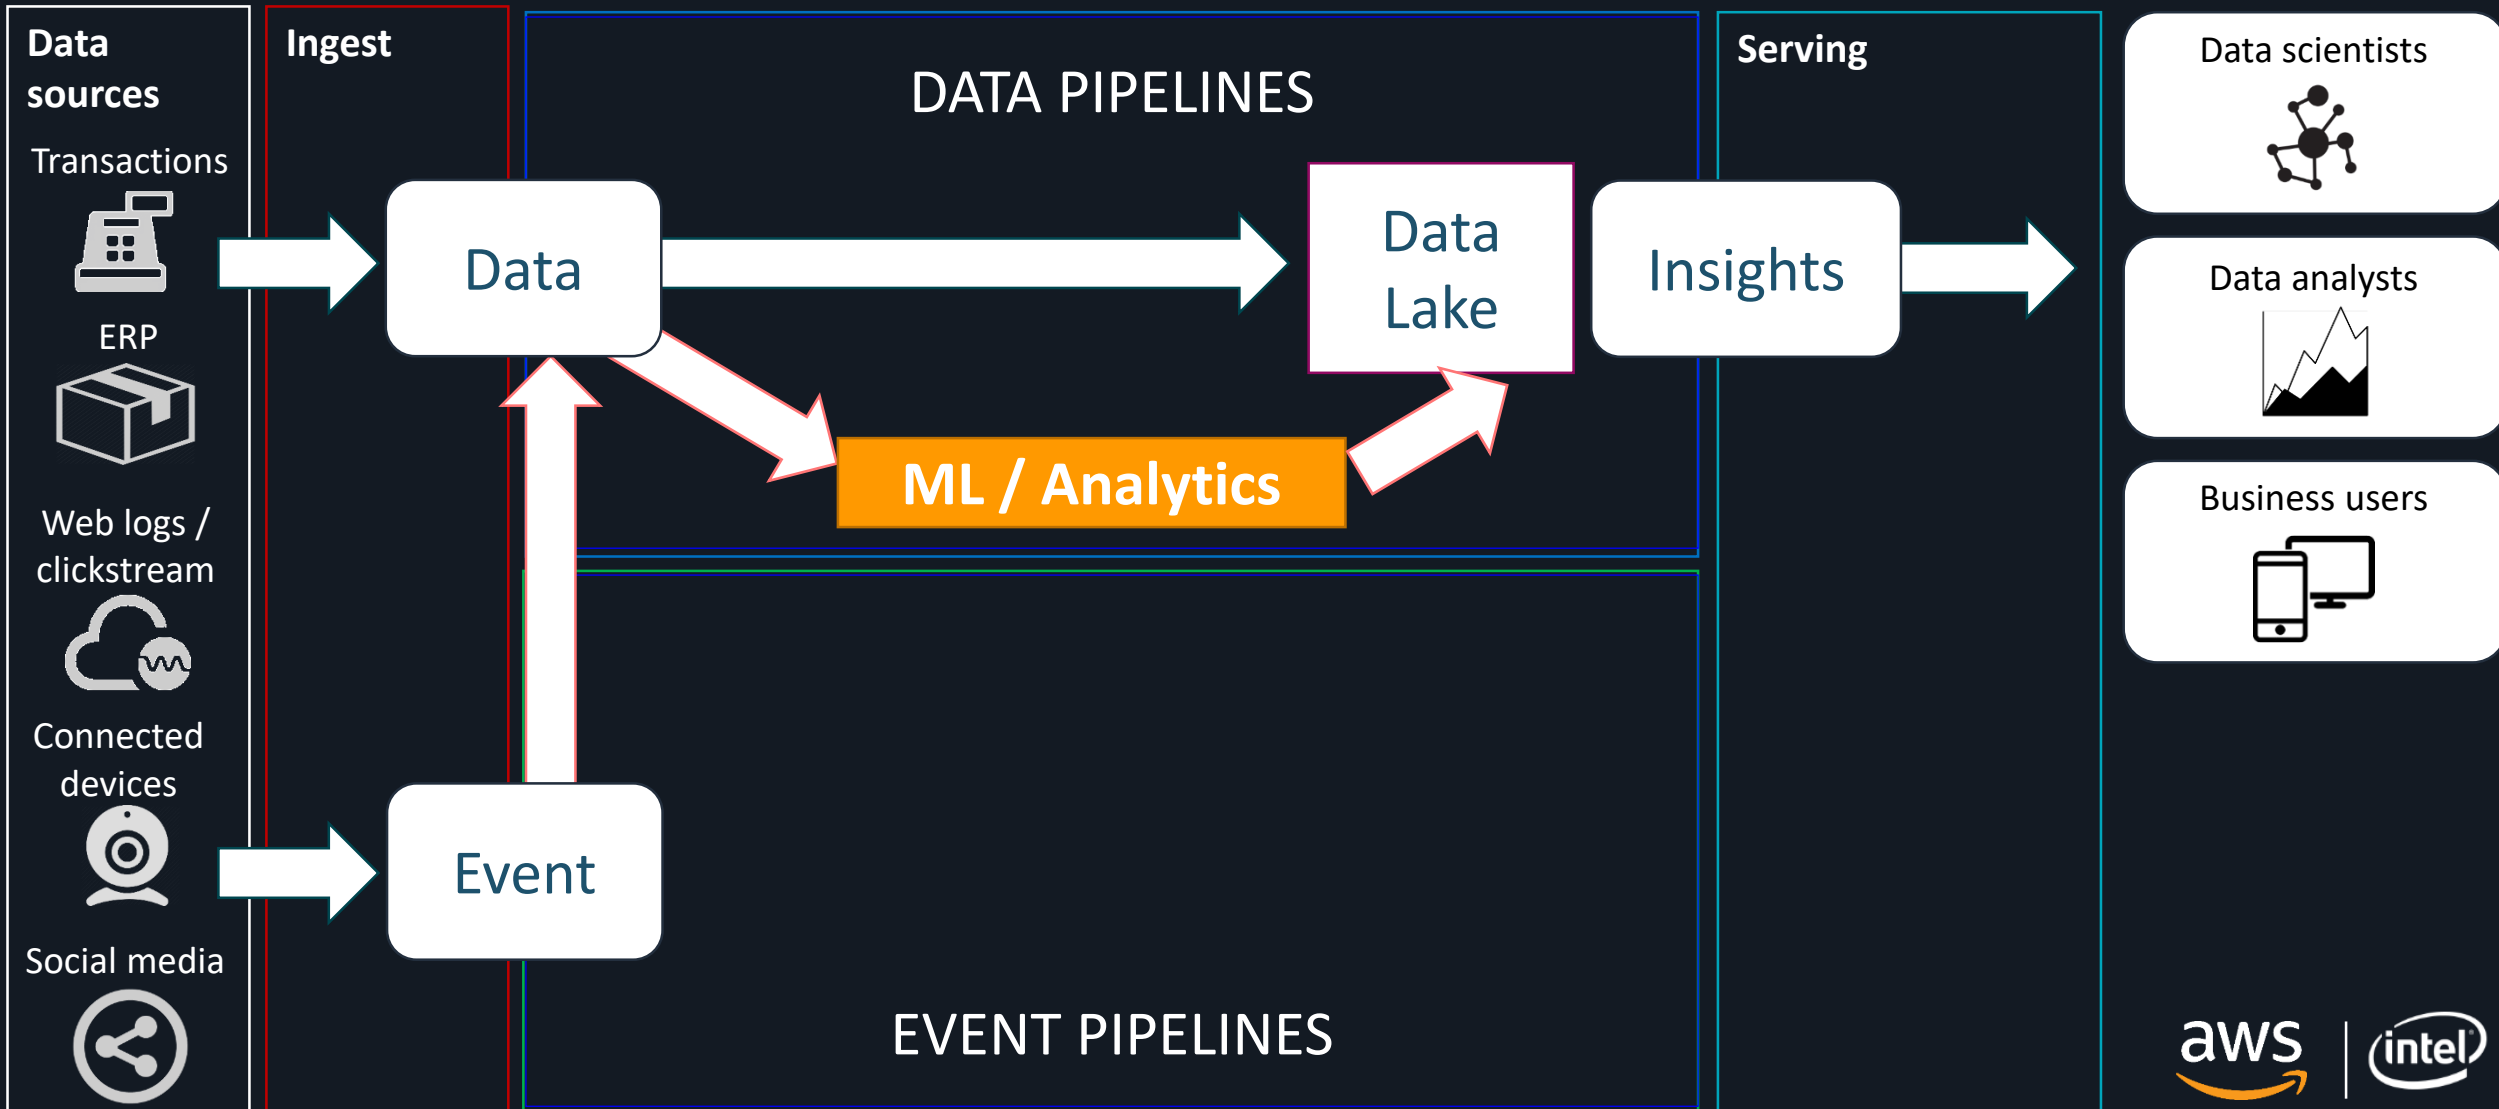
Insights to enhance business applications, new digital services





# Modern data architecture

Innovate for new revenues - personalization and forecasting



# Outcome 3: Real-time engagement

# EarEcstasy offers a personalized life soundtrack

Personalized, based on  
past preferences,  
people with similar behaviors,  
and environments detected



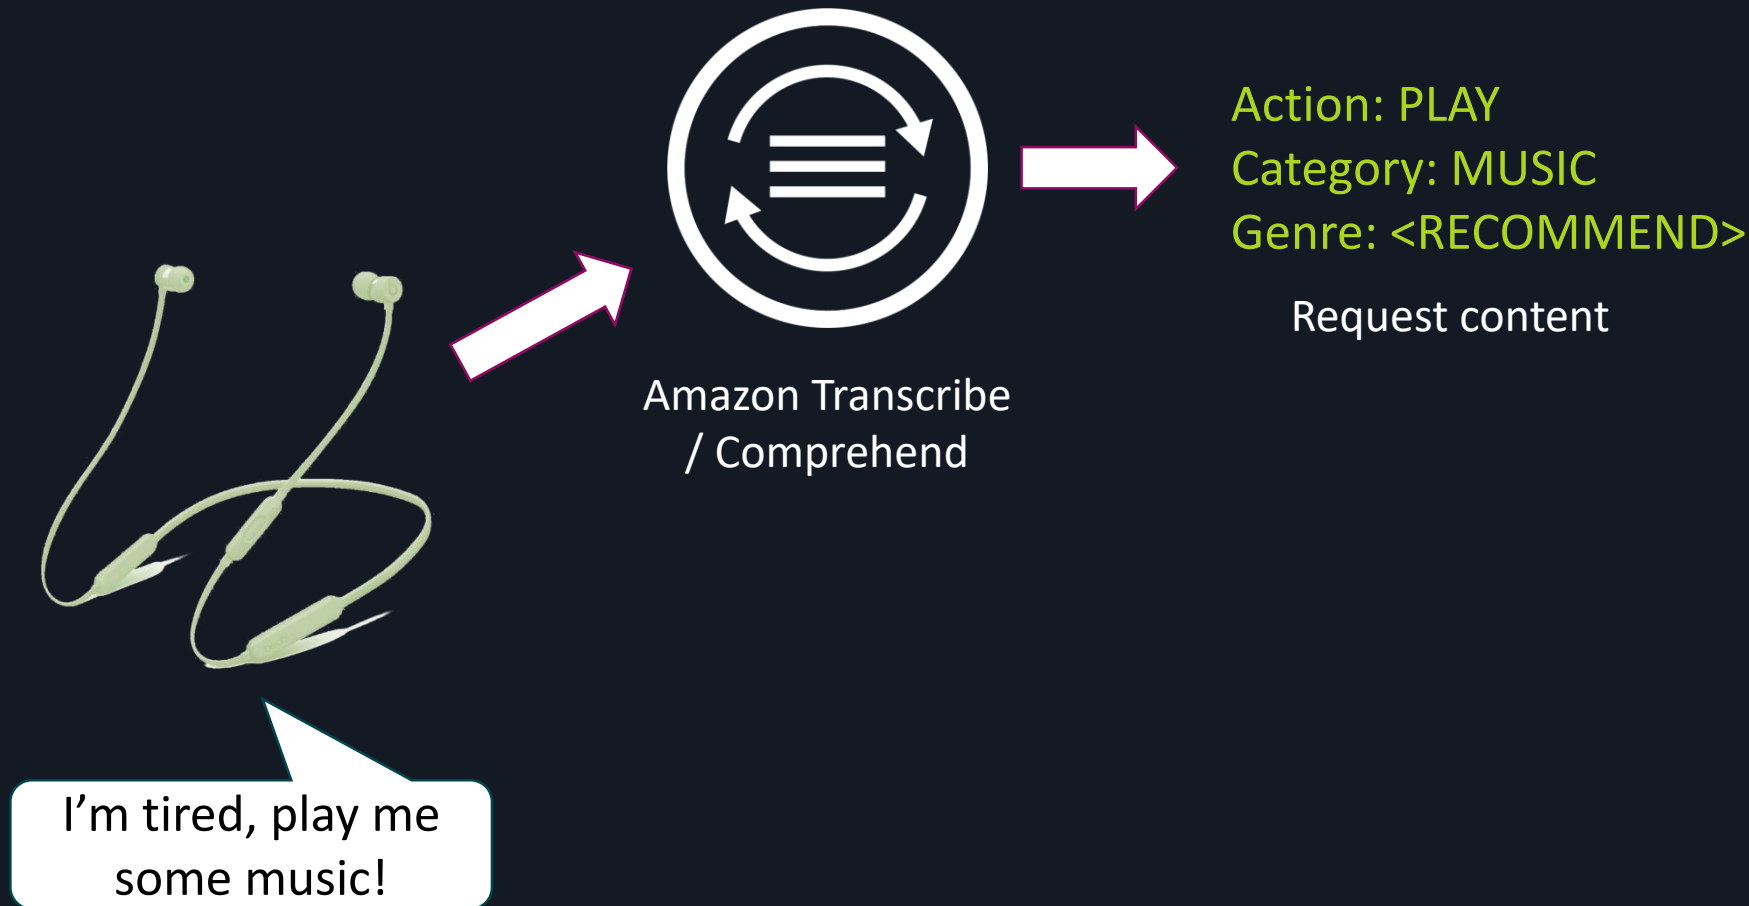
# Use EarEcstasy voice enablement to play music

# Use EarEcstasy voice enablement to play music

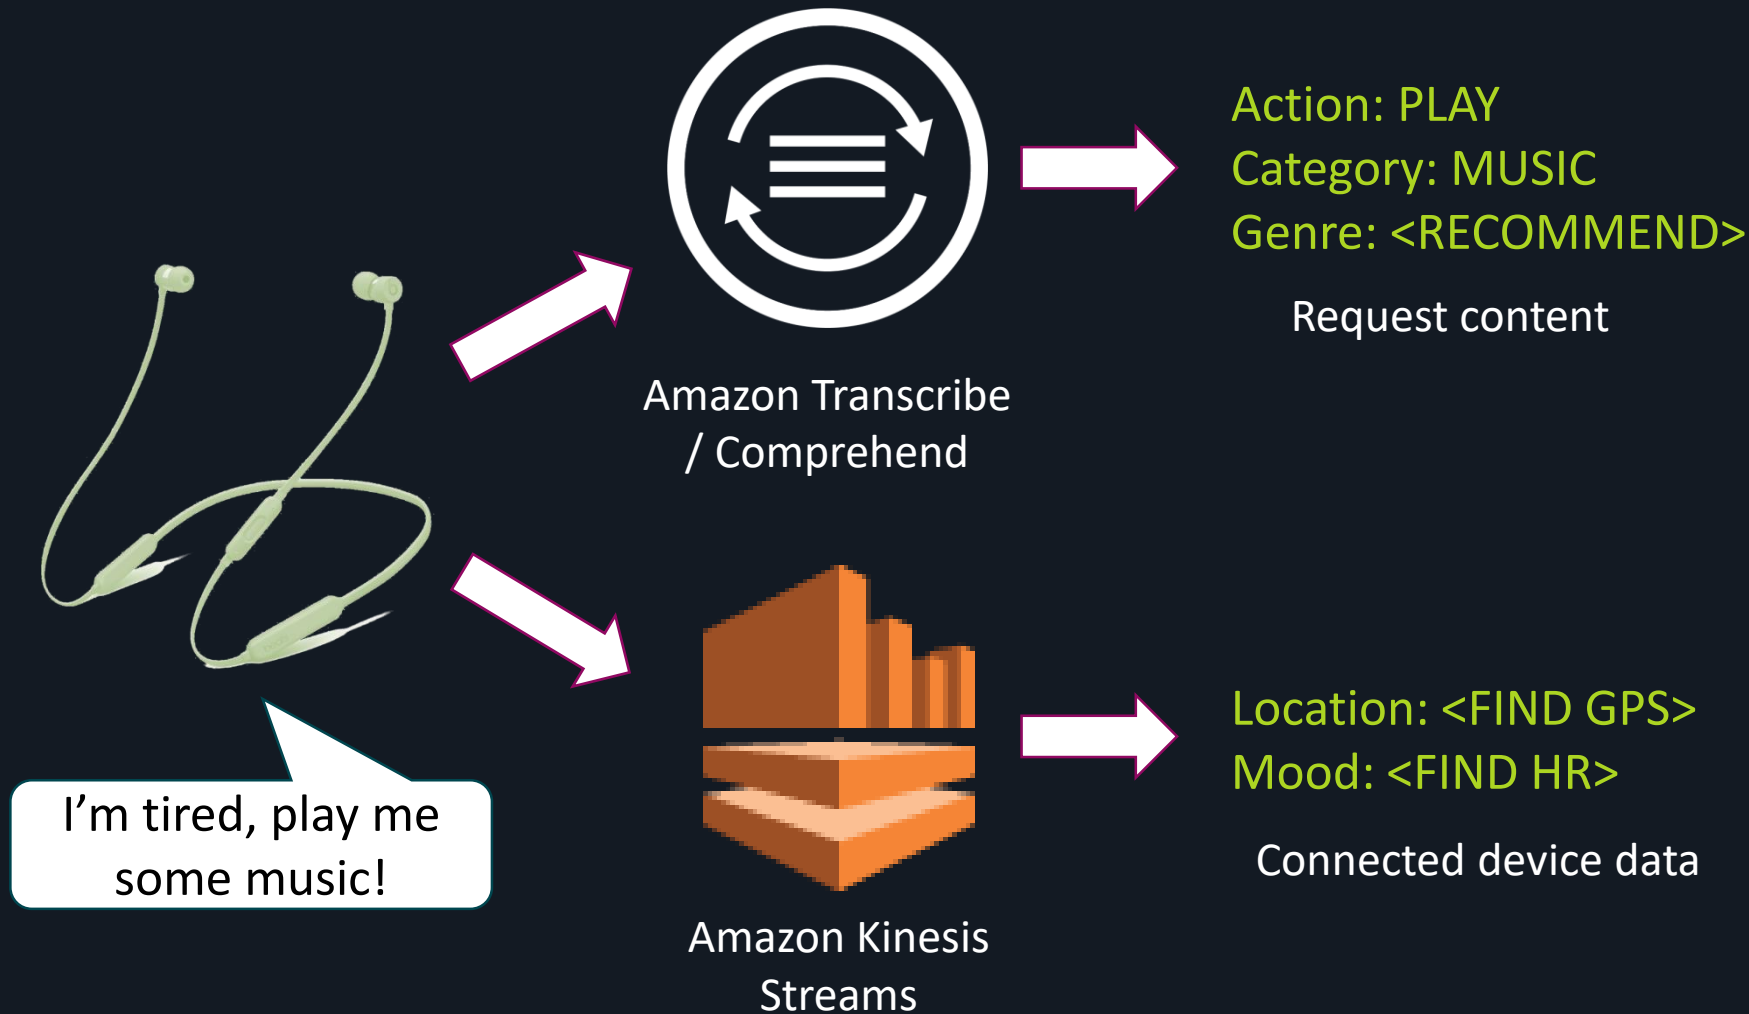


I'm tired, play me  
some music!

# Use EarEcstasy voice enablement to play music



# Use EarEcstasy voice enablement to play music

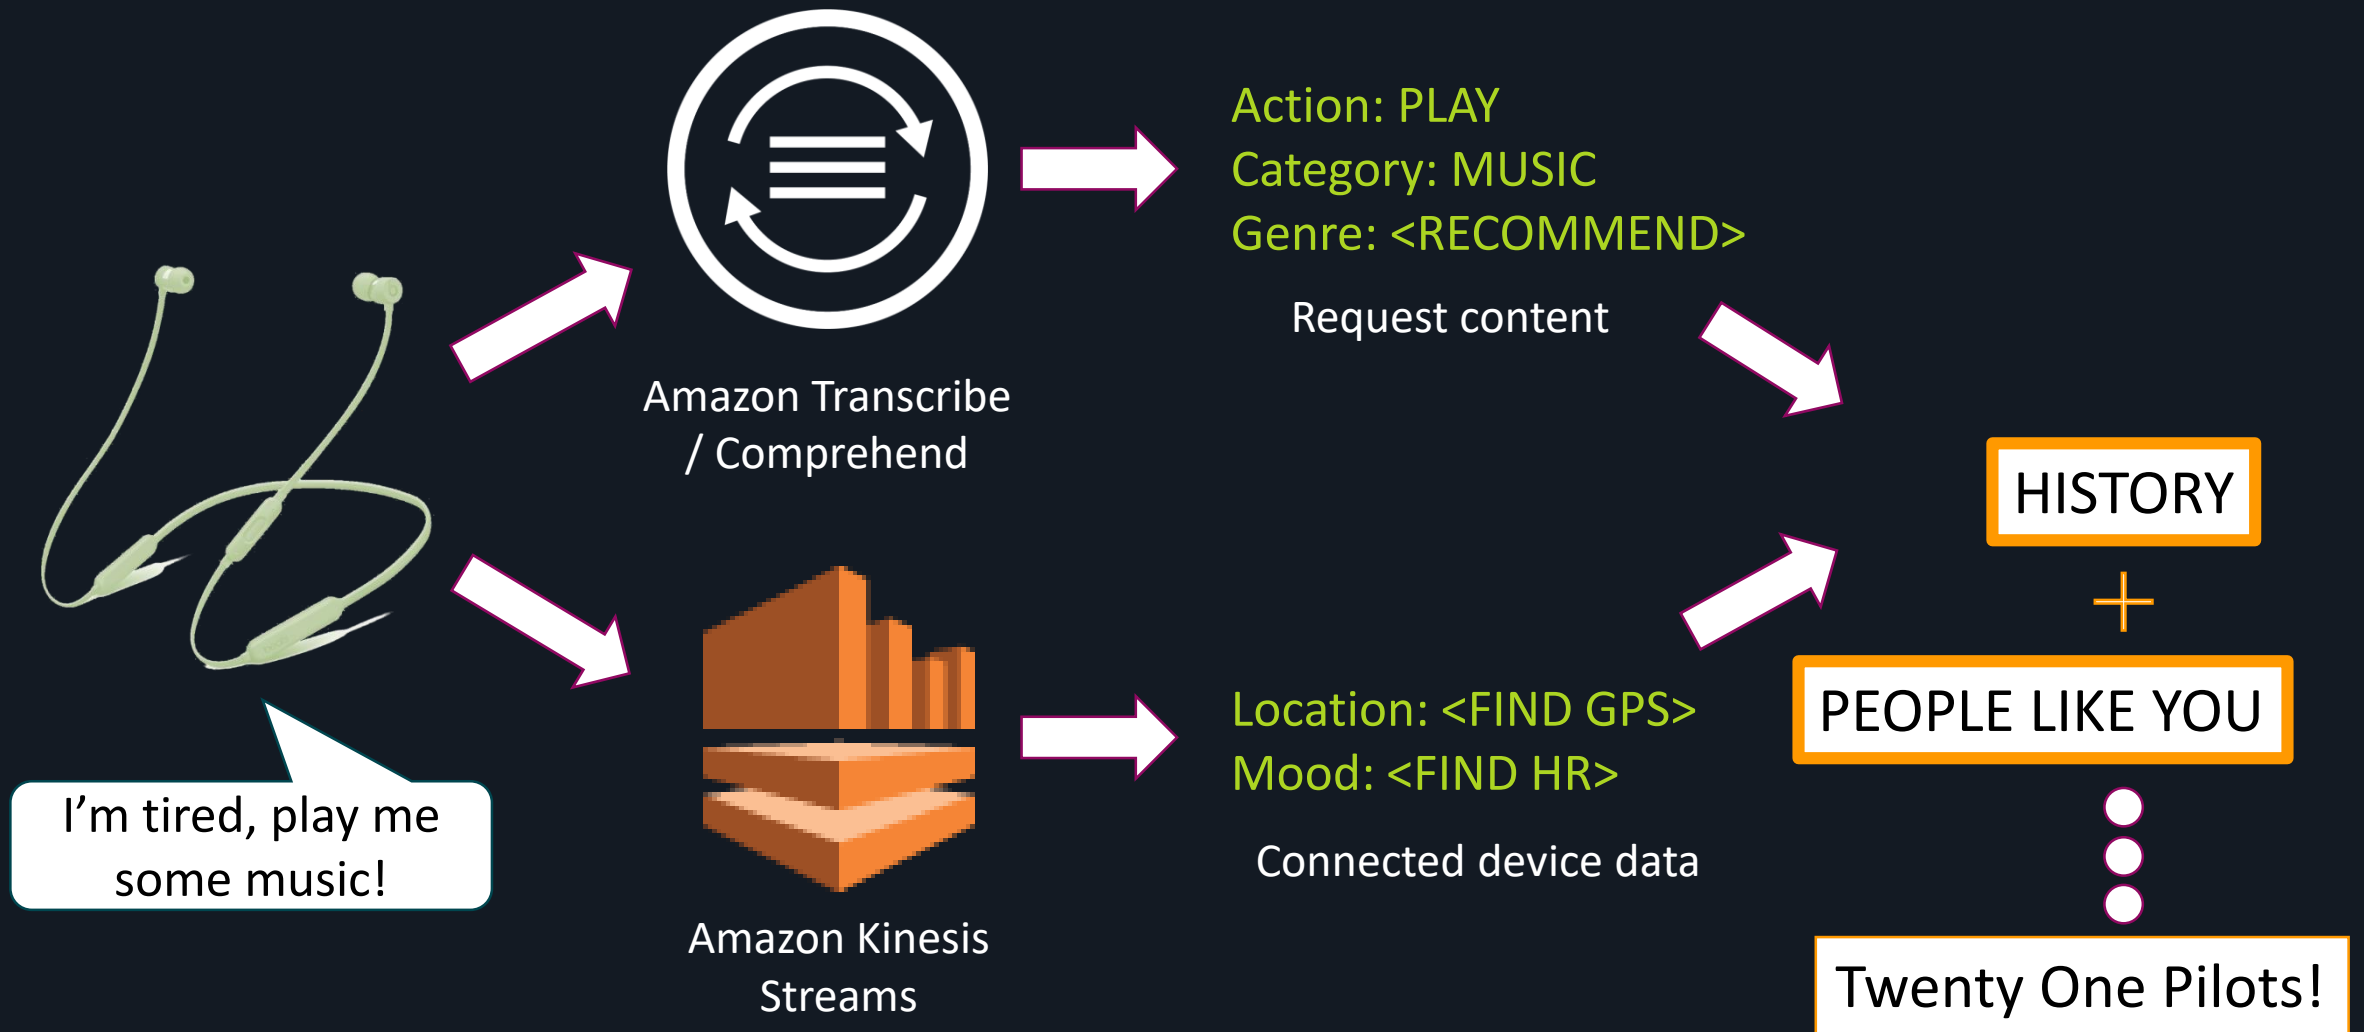


# Use EarEcstasy voice enablement to play music






# Use EarEcstasy voice enablement to play music



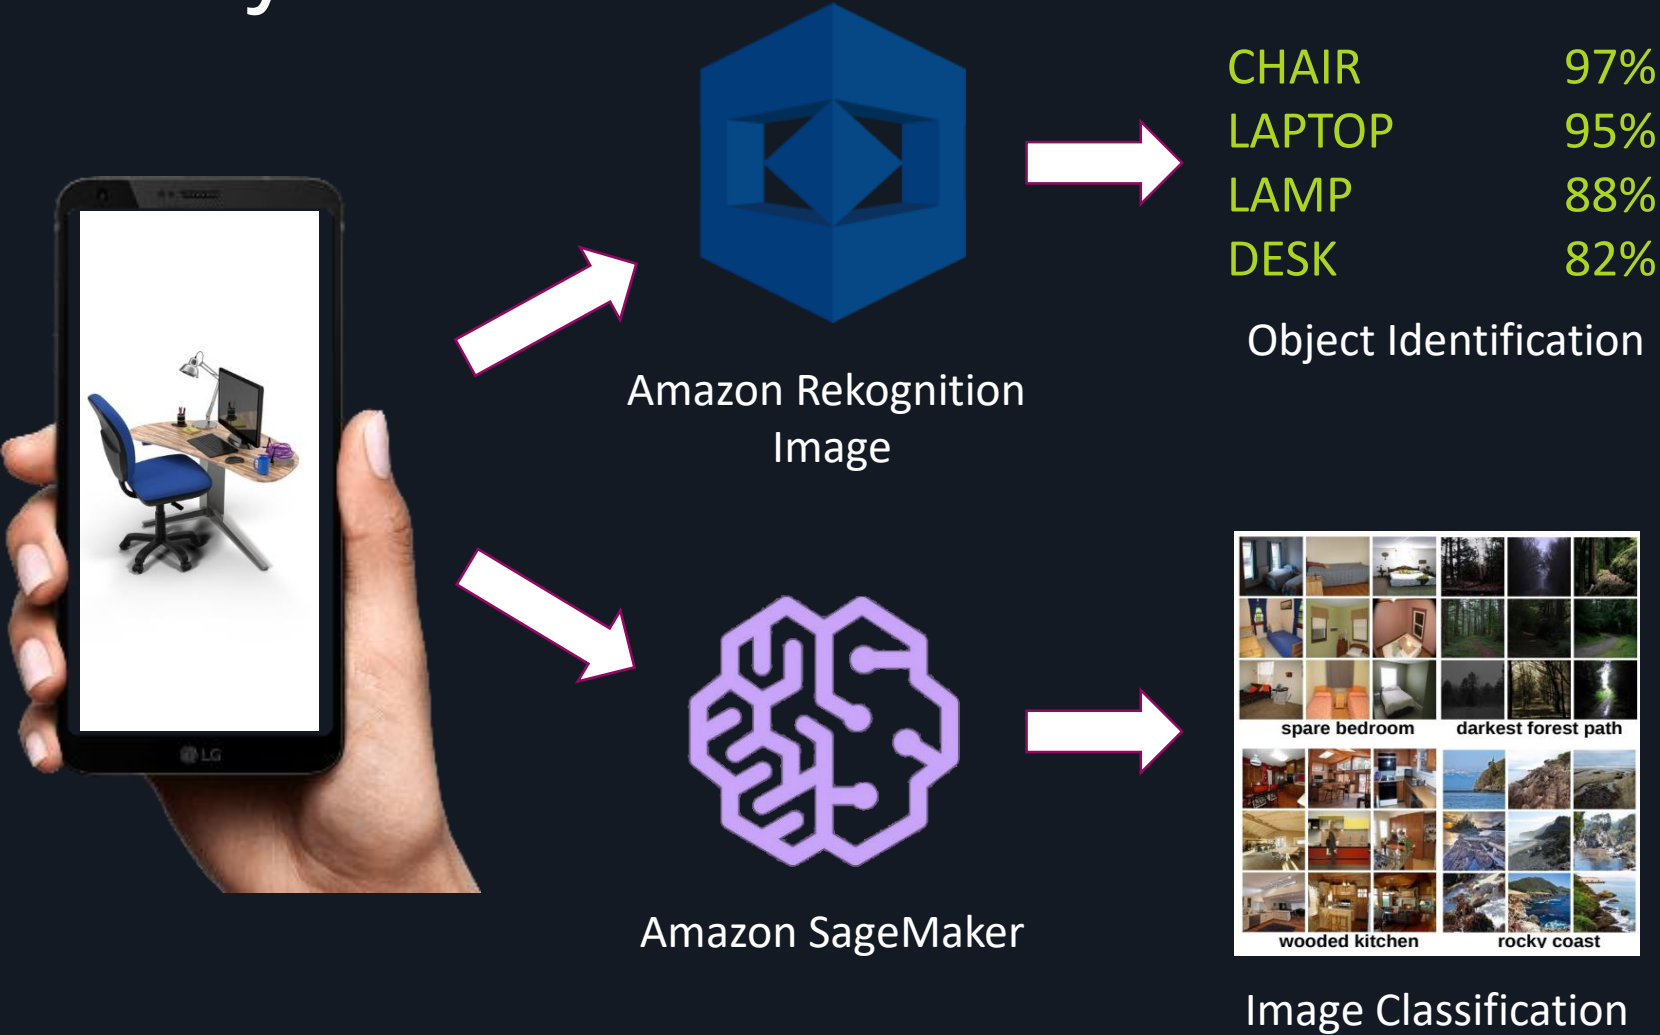
# Use the mobile app to take a picture to identify activity



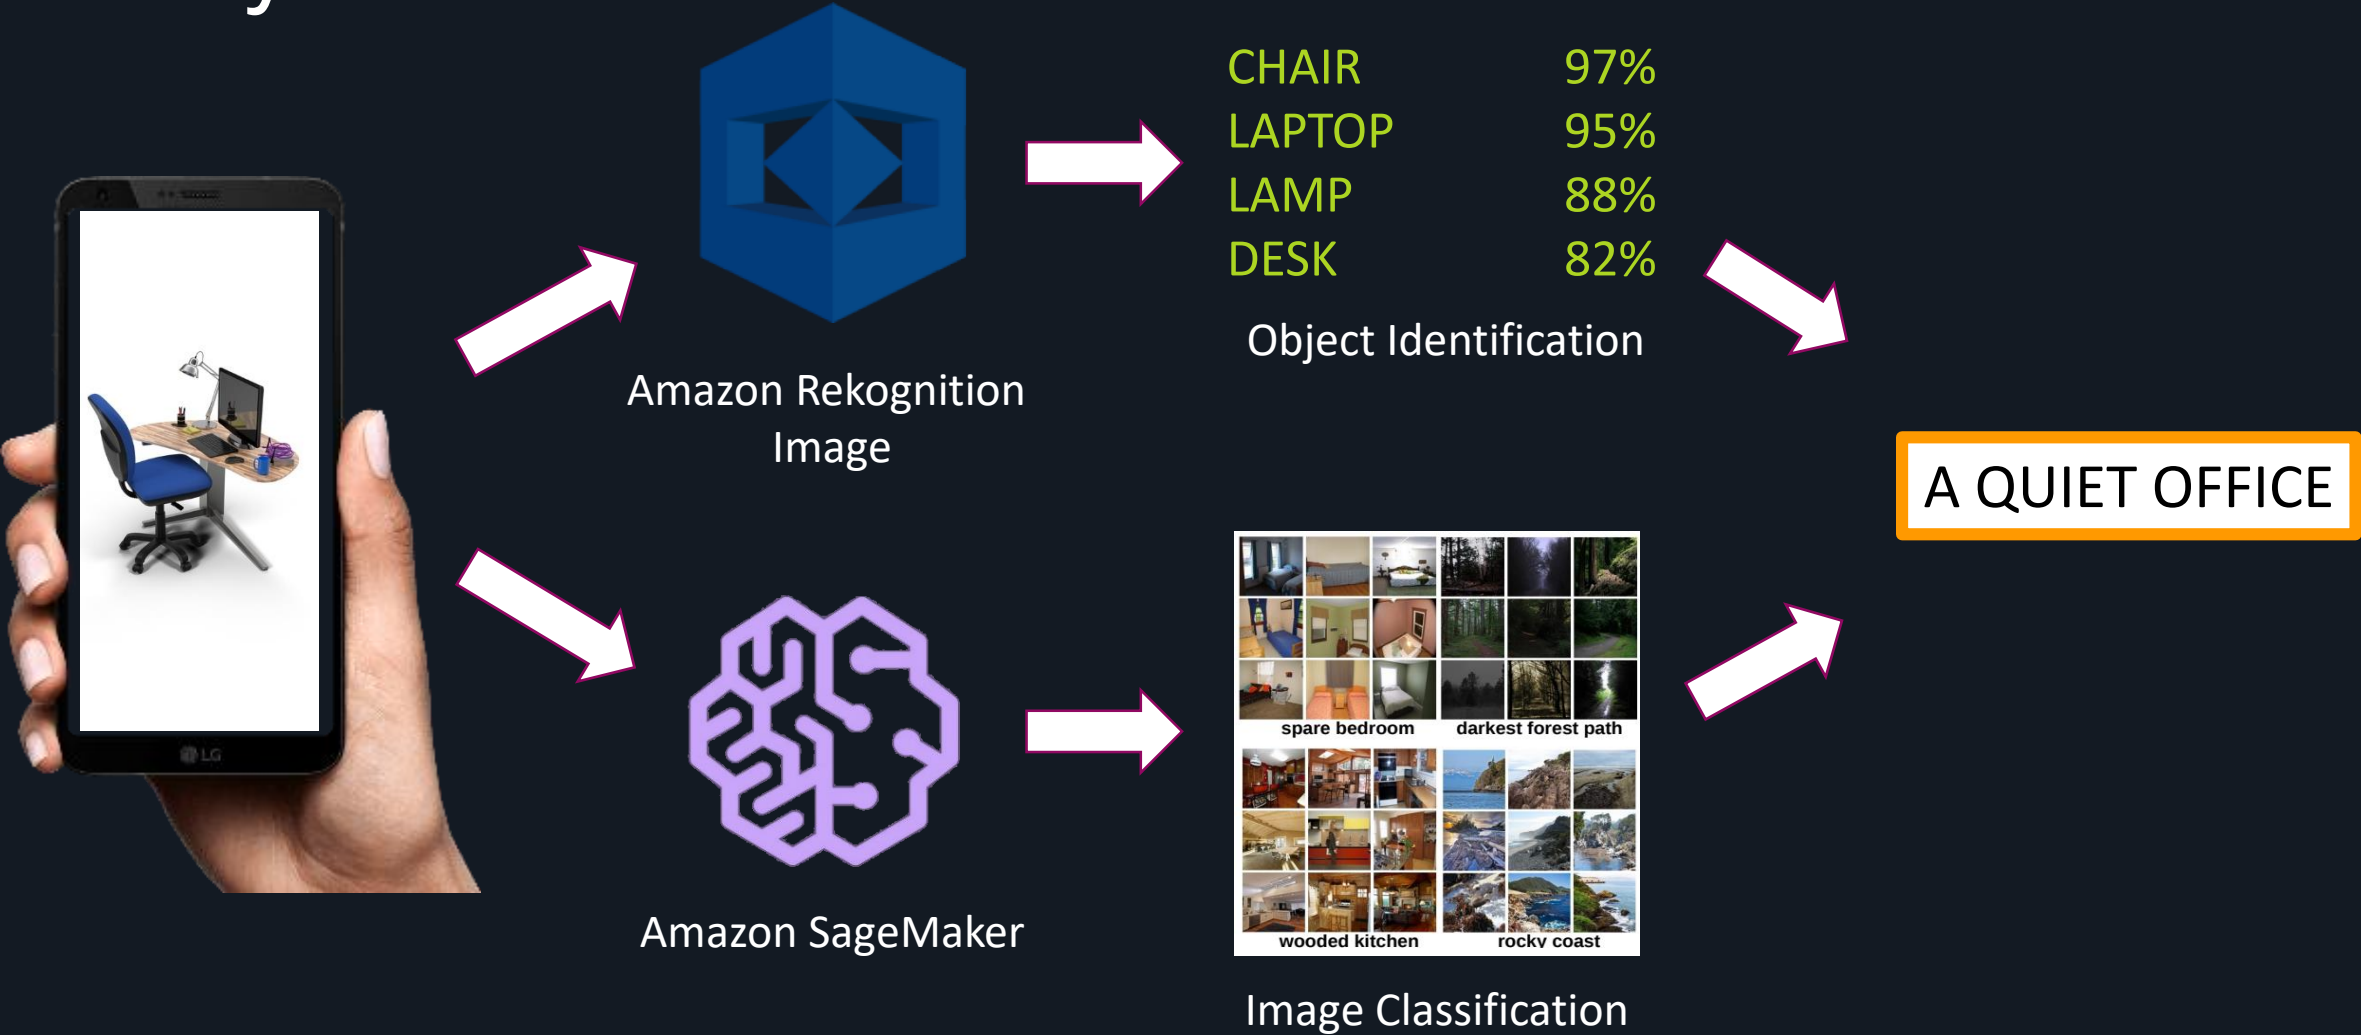
# Use the mobile app to take a picture to identify activity



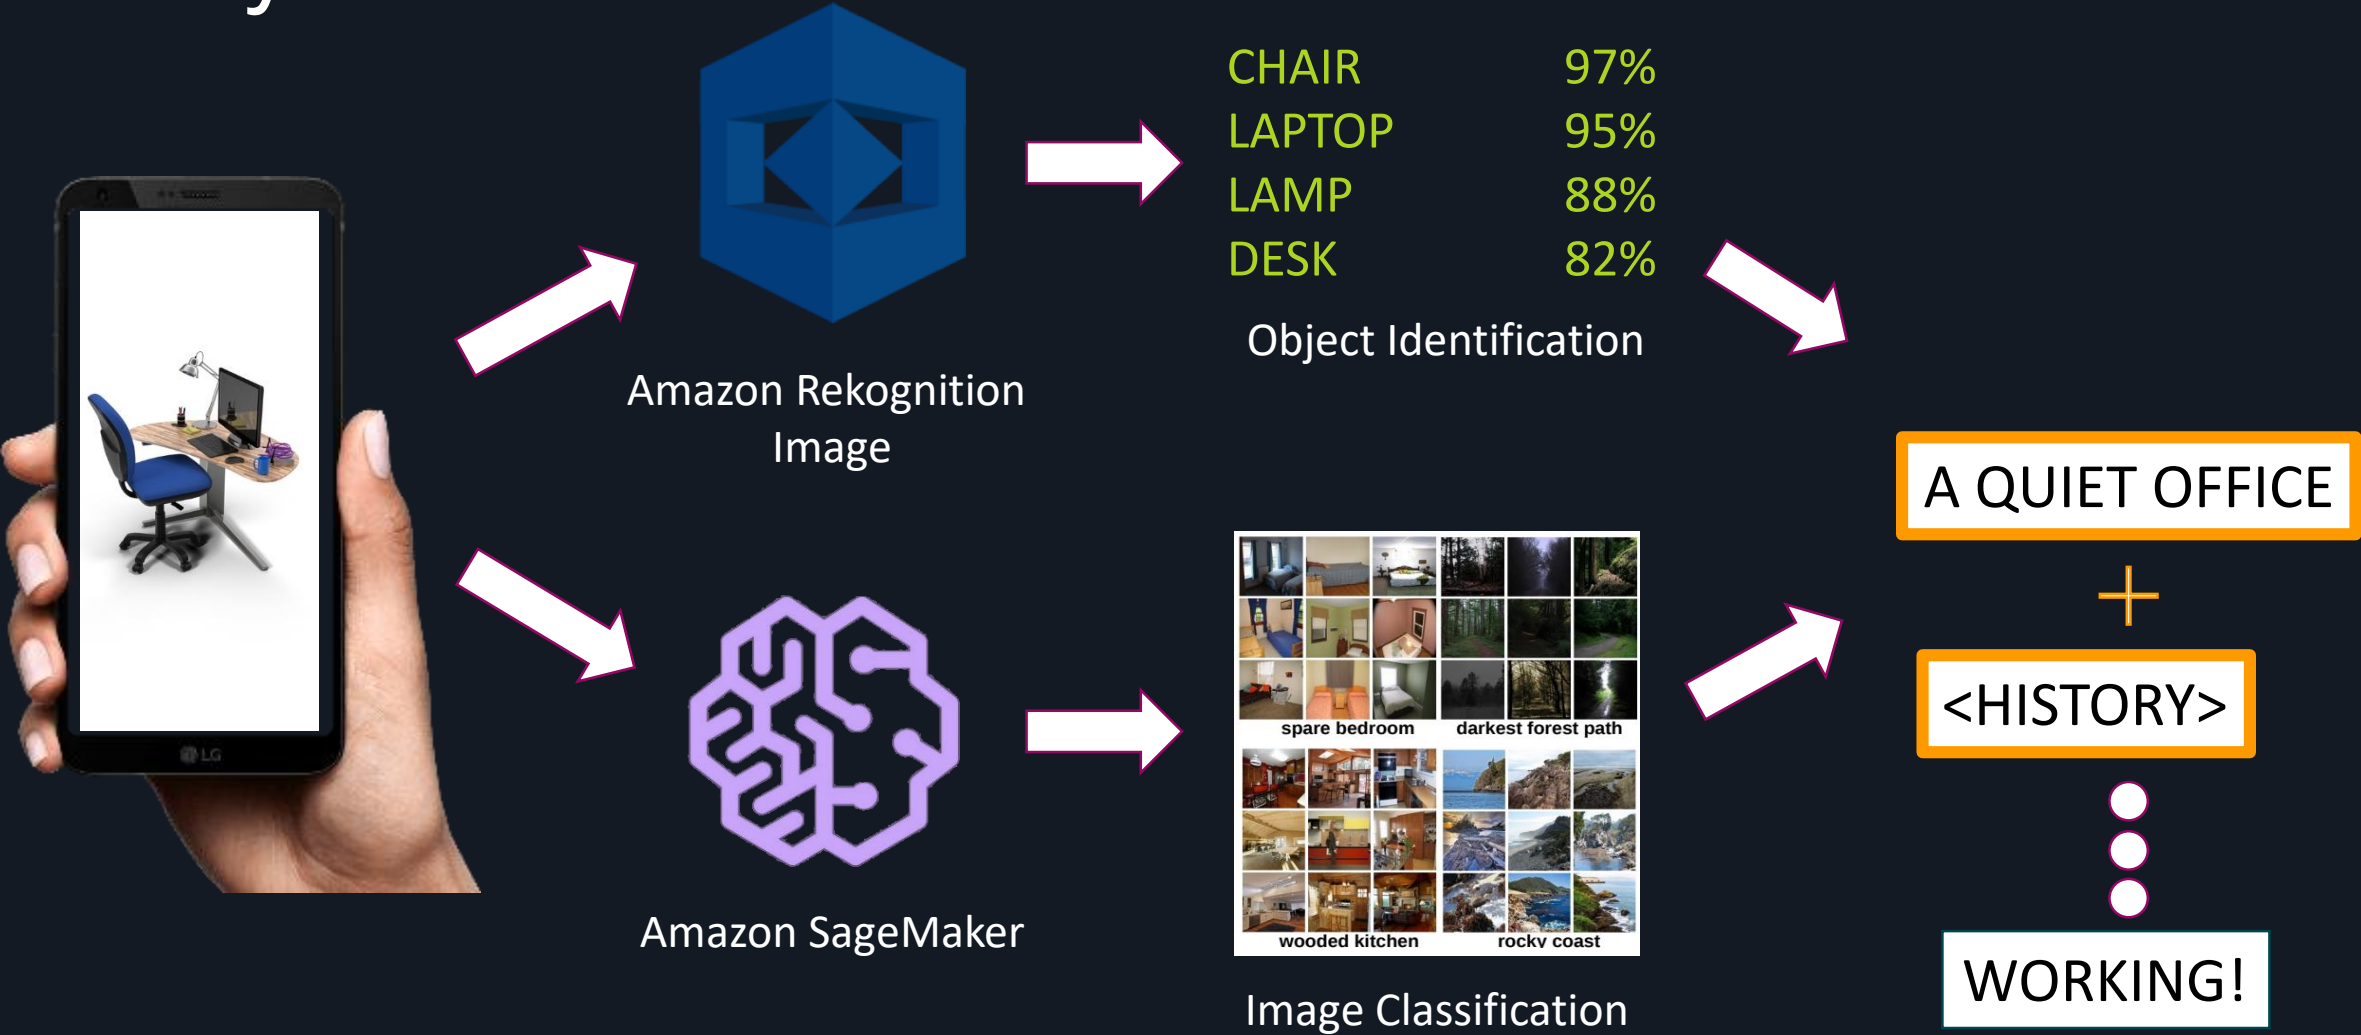
# Use the mobile app to take a picture to identify activity



# Use the mobile app to take a picture to identify activity

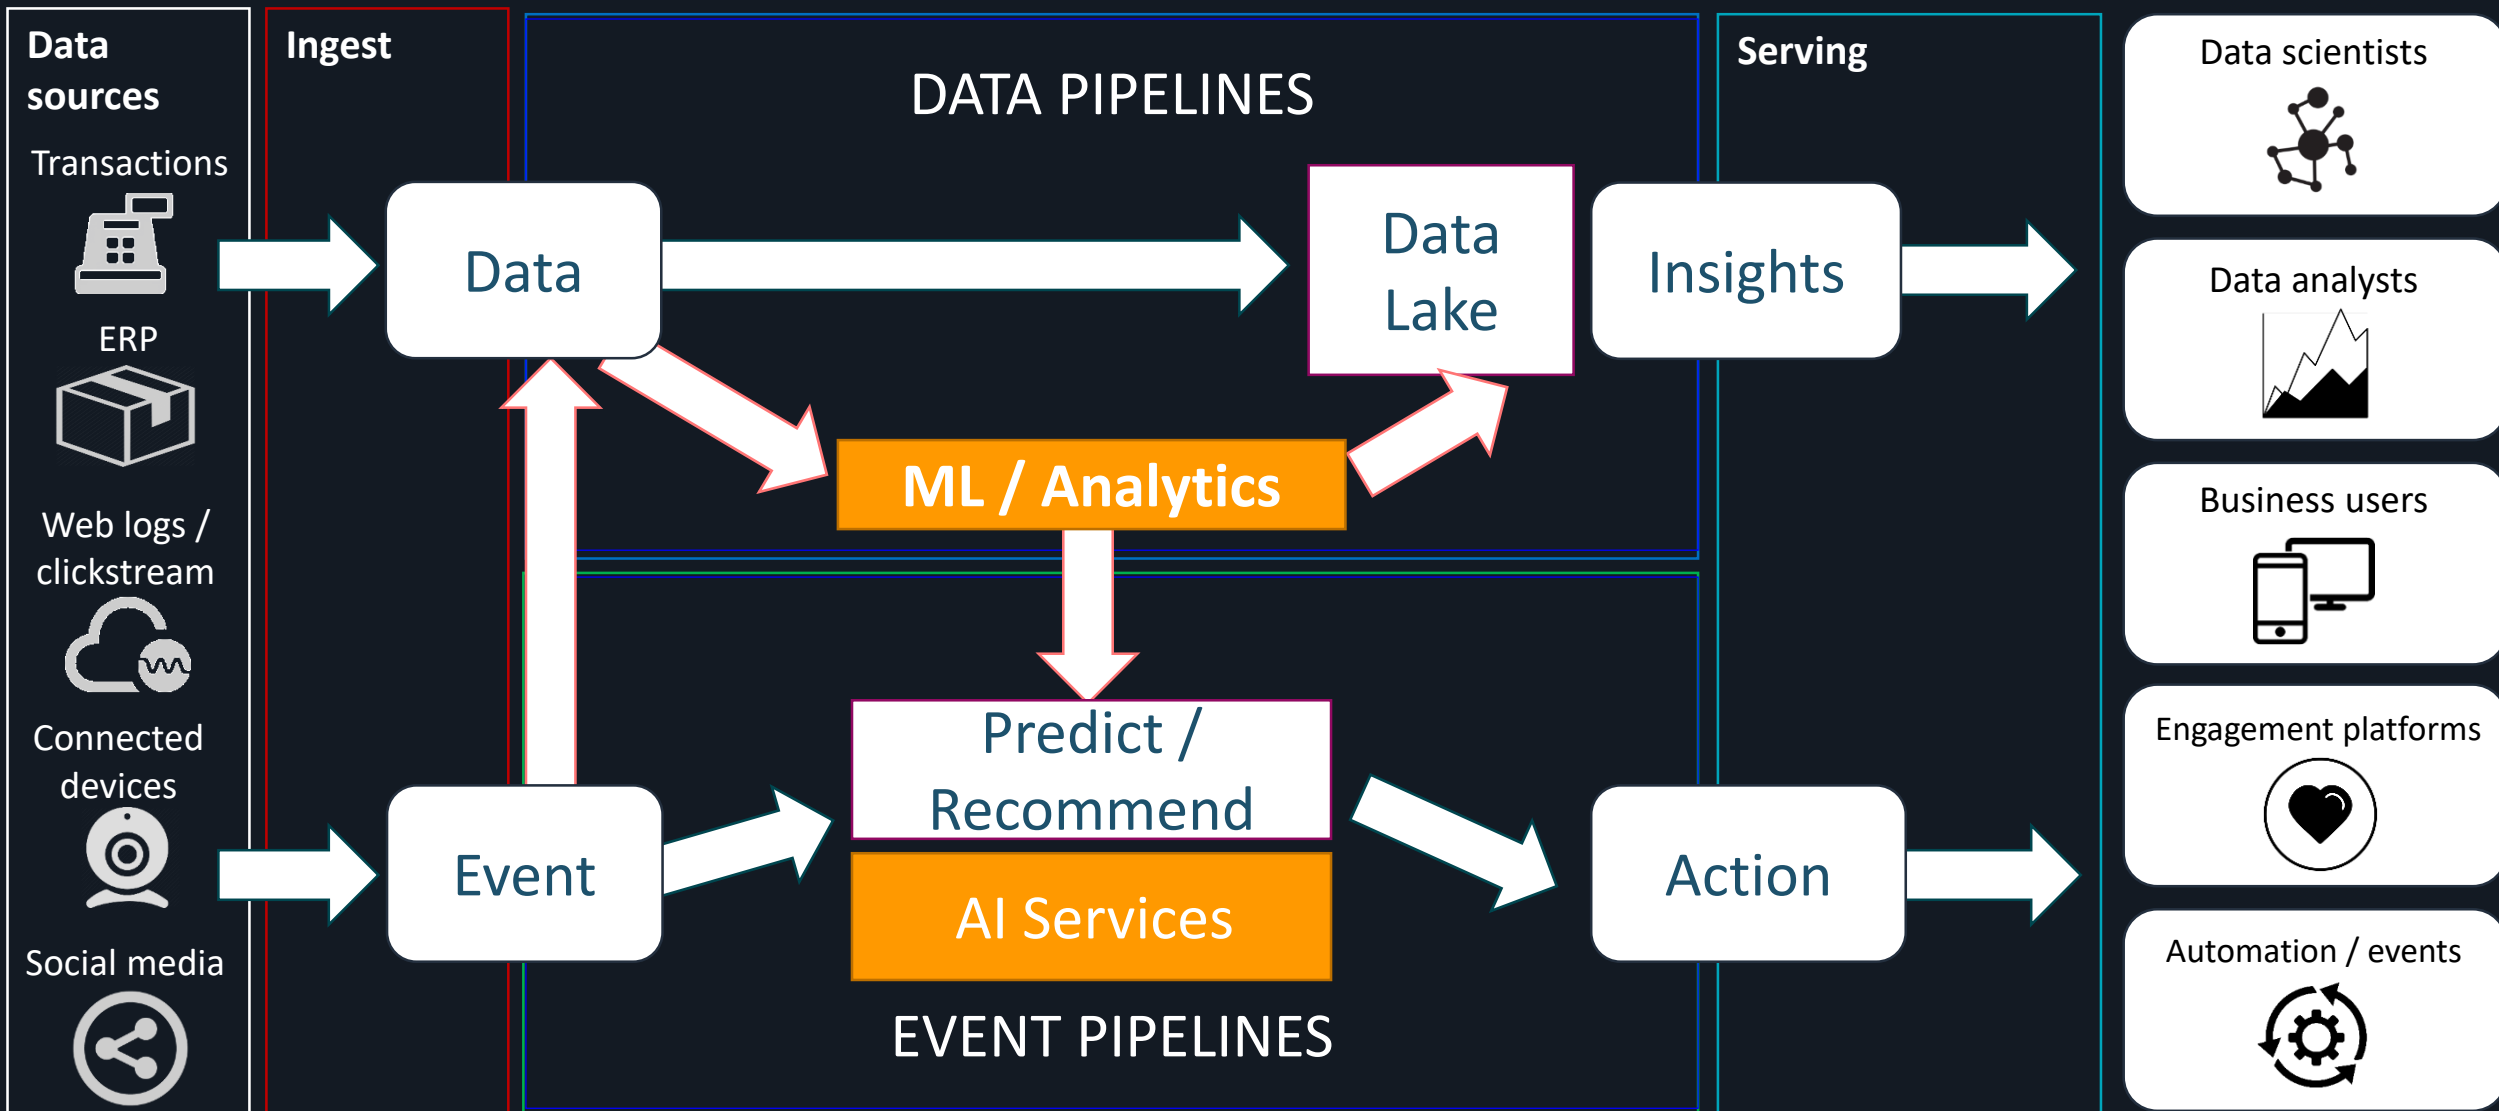


# Use the mobile app to take a picture to identify activity



# Modern data architecture

Real-time engagement and interactive customer experiences



# Business Outcomes on a Modern Data Architecture



# Business Outcomes on a Modern Data Architecture

## Outcome 1 : Modernize and consolidate

Insights to enhance business applications and create new digital services

# Business Outcomes on a Modern Data Architecture

## Outcome 1 : Modernize and consolidate

Insights to enhance business applications and create new digital services

## Outcome 2 : Innovate for new revenues

Personalization, demand forecasting, risk analysis

# Business Outcomes on a Modern Data Architecture

## Outcome 1 : Modernize and consolidate

Insights to enhance business applications and create new digital services

## Outcome 2 : Innovate for new revenues

Personalization, demand forecasting, risk analysis

## Outcome 3 : Real-time engagement

Interactive customer experience, event-driven automation, fraud detection

# Ready to build better business from your ideas?

# Ready to build better business from your ideas?



Short list projects that  
directly impact customer  
engagement and adoption

# Ready to build better business from your ideas?



Short list projects that directly impact customer engagement and adoption

Build simple data pipelines that allow you to test new ideas, and fill your data lake

# Ready to build better business from your ideas?



Short list projects that directly impact customer engagement and adoption



Build simple data pipelines that allow you to test new ideas, and fill your data lake



Ask our solution architects and professional services teams to help you build

# Learn from AWS experts. Advance your skills and knowledge. Build your future in the AWS Cloud.



## Digital Training

Free, self-paced online courses built by AWS experts



## Classroom Training

Classes taught by accredited AWS instructors



## AWS Certification

Exams to validate expertise with an industry-recognized credential

Ready to begin building your cloud skills?  
Get started at: <https://www.aws.training/>



# With deep expertise on AWS, APN Partners can help your organization at any stage of your Cloud Adoption Journey.



## **AWS Managed Service Providers**

APN Consulting Partners who are skilled at cloud infrastructure and application migration, and offer proactive management of their customer's environment.




## **AWS Competency Partners**

APN Partners who have demonstrated technical proficiency and proven customer success in specialized solution areas.



## **AWS Marketplace**

A digital catalog with thousands of software listings from independent software vendors that make it easy to find, test, buy, and deploy software that runs on AWS.



## **AWS Service Delivery Partners**

APN Partners with a track record of delivering specific AWS services to customers.

**Ready to get started with an APN Partner?**  
**Find a partner: <https://aws.amazon.com/partners/find/>**  
**Learn more at the AWS Partner Network Booth**

Track	Migration and Workloads	Database and Storage	Big Data and Analytics	AI and Machine Learning	IoT	Build on AWS I	Build on AWS II	Building Practices with AWS	Build Your Own Architecture	Live Coding
30 mins	Migrate Away That Technical Debt (Level 200)	AWS Purpose-Built Database Strategy: The Right Tool for The Right Job (Level 200)	Build Business Outcomes with a Modern Data Architecture (Level 200)	Demystifying Machine Learning on AWS (Level 200)	Drive Successful Business Outcomes with IoT on AWS (Level 200)	Streamline Your Serverless Development with AWS CodeStar and AWS Cloud9 (Level 200)	AWS Security for Builders: Best Practices (Level 200)	TechShift: Migrating and Replatforming Your Software on AWS (Level 200)	Step by Step Guide to Build Your Own Video On-Demand Architecture (Level 200)	Machine Learning in Production with Twitter Bots (Level 400)
30 mins	Optimize Cost and Efficiency on AWS: Tips and Tools (Level 200)	Build High-Performance Apps with In-Memory Data (Level 200)	Serving Business Insights with the AWS Analytics Suite (Level 200)	Detect Anomalies in Your Data with Amazon SageMaker (Level 300)	Connecting Devices to AWS IoT Core (Level 300)	Run Containers without Managing Servers or Clusters with AWS Fargate (Level 300)	Improving Release Velocity with Continuous Delivery on AWS (Level 300)	The 4 Secrets of Successful SaaS Companies (Level 200)	1. Ingest Your Data 2. Scheduling and Transcoding Data 3. Content Delivery and Scaling	
30 mins	SAP on AWS: Big Businesses, Big Workloads, Big Data (Level 200)	Supercharge Your Apps with Amazon Neptune Graph Database (Level 200)	Designing Data Lakes: Best Practices (Level 200)	Extract Data from Images and Videos with Amazon Rekognition (Level 300)	Leveraging IoT at the Edge and Industrial Workloads (Level 300)	Comparing Container Options for Microservices on AWS (Level 200)	Taking Your Progressive Web App to the Next Level With AWS AppSync and GraphQL (Level 300)	Building Scalable Migration Practices for Enterprise Applications (Level 200)	4. Search and Recommendations Using Machine Learning Services 5. Bring Smarts to Your Media Player Using Real-Time Analytics	
30 mins	Migrate Microsoft Applications to AWS like an Expert (Level 200)	Migrate Large-Scale Datasets to AWS (Level 200)	Data Pipelines with AWS Glue (Level 200)	Manage Deep Learning Models at Scale with Amazon SageMaker (Level 300)	Managing Complexities of IoT Workloads at Scale (Level 300)	Self-Defending Borders: A Developer's Approach to Security (Level 300)	Build an Omni-Channel Experience with Amazon Connect and Amazon Lex (Level 200)	Building a Big Data and Analytics Practice: From Zero to Hero in 5 Steps (Level 200)	6. Best Practices on Going Global	Build Serverless Data-Driven Web Apps Using React, GraphQL, and AWS AppSync (Level 400)
30 mins	VMware Cloud on AWS: Next Generation Hybrid Cloud Architecture (Level 200)	Protect Your Data Assets with AWS (Level 200)	Enterprise Data Warehouse Offload and Migration (Level 200)	Better Customer Experiences with Chatbots and Conversational Interfaces (Level 300)	Putting It All Together: Building an IoT Solution (Level 200)	Deliver Software Quickly and Safety with Serverless Applications (Level 300)	Build High Quality Mobile Apps in Minutes with AWS (Level 200)	Build a Next-Generation MSP Practice with AWS (Level 100)	Succeed in the Cloud with AWS Training (Level 100)	

# Thank You for Attending AWS Innovate

We hope you found it interesting! A kind reminder to **complete the survey.**

Let us know what you thought of today's event and how we can improve the event experience for you in the future.



[aws-apac-marketing@amazon.com](mailto:aws-apac-marketing@amazon.com)



[twitter.com/AWSCloud](https://twitter.com/AWSCloud)



[facebook.com/AmazonWebServices](https://facebook.com/AmazonWebServices)



[youtube.com/user/AmazonWebServices](https://youtube.com/user/AmazonWebServices)



[slideshare.net/AmazonWebServices](https://slideshare.net/AmazonWebServices)



[twitch.tv/aws](https://twitch.tv/aws)